

Time Series Representations with Hard-Coded Invariances

Anonymous Authors¹

Abstract

Automatically extracting robust representations from large and complex time series data is becoming imperative for several real-world applications. Unfortunately, the potential of common neural network architectures in capturing invariant properties of time series remains relatively underexplored. For instance, convolutional layers often fail to capture underlying patterns in time series inputs that encompass strong deformations, such as trends. Indeed, invariances to some deformations may be critical for solving complex time series tasks, such as classification, while guaranteeing good generalization performance. To address these challenges, we mathematically formulate and technically design efficient and hard-coded *invariant convolutions* for specific group actions applicable to the case of time series. We construct these convolutions by considering specific sets of deformations commonly observed in time series, including *scaling*, *offset shift*, and *trend*. We further combine the proposed invariant convolutions with standard convolutions in single embedding layers, and we showcase the layer capacity to capture complex invariant time series properties in several scenarios.

1. Introduction

Recently, there has been a growing interest in applying machine learning to time series data, further necessitated by the challenging properties of time series, including varying modalities, high noise levels, and distribution shifts.

Machine learning for time series has gradually moved from statistical methods, e.g., autoregressive models (Box et al., 2015) and dictionary-based classification (Middlehurst et al., 2019), to neural networks. Notable architectures are built upon recurrent layers, convolutional layers, and, more re-

cently, transformers. More specifically, convolutional neural networks (CNNs) consistently lead time series tasks such as classification and clustering (Ismail Fawaz et al., 2019; Tonekaboni et al., 2021). Often integrated with other modules, CNNs excel in feature extraction while offering interpretability through kernel weight visualization.

Advances in invariance modeling for deep learning have been achieved with proper mathematical formalism of group action on data, notably images and graphs (Kvinge et al., 2022; Bronstein et al., 2021). Two strategies for incorporating invariance in deep networks are learning them through data augmentation, such as contrastive learning (Antoniou, 2017) or hard-coding invariances directly within the network architecture. Translation in CNNs and permutation in graph neural networks (GNNs) are two prominent examples of hard-coded invariances in terms of architectural design, among others (Bietti & Mairal, 2019; Horie et al., 2020). Relevant works range across sets, images, point clouds, and graphs (Zaheer et al., 2017; Keriven & Peyré, 2019).

Incorporating knowledge on time series invariances during training is an emerging field of study. Several works focus on learning invariances, e.g., contrastive learning that leverages augmentations and customizable losses (Franceschi et al., 2019; Eldele et al., 2021). Yet, the selection of views within the time series collection to contrast, as well as the types of transformations to consider (e.g., scaling, shifting), is often arguable within the research community (Yue et al., 2022). While contrastive learning constitutes an implicit way of introducing invariances into learning, very few studies leverage time series hard-coded invariances. In this work, we aim to embed invariances into the network design, similar to approaches that achieve permutation invariance in graphs (Maron et al., 2018) and local translation and rotation invariance in images (Weiler & Cesa, 2019). Such invariant networks offer improved generalization properties compared to their learned counterparts at the expense of additional computational costs (Kvinge et al., 2022).

Enforcing deformation-specific invariances in network design can significantly improve performance for time series tasks. For example, the trend in time series constitutes a common deformation, and removing its effects is an active research area known as baseline removal (Hipke et al., 2019; Zhang et al., 2010). Modern time series machine

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

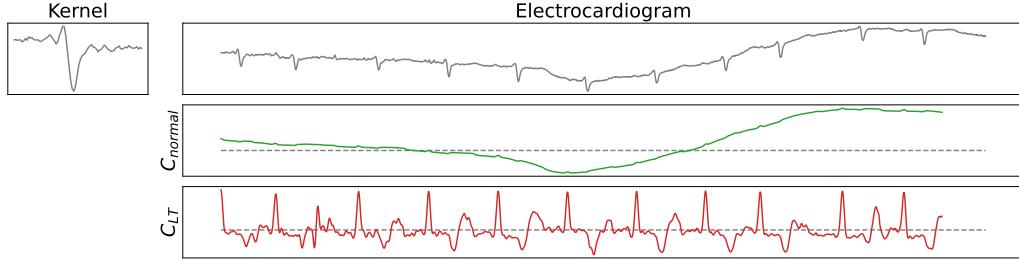


Figure 1: **Top right:** Segment of a electrocardiogram (ECG) from the MIT-BIH dataset (Goldberger et al., 2000; Moody & Mark, 2001). **Top left:** The convolutional kernel is the first heartbeat. **Middle:** With the normal convolution, individual heartbeats are not identifiable as they are blurred by the trend. **Bottom:** With linear trend invariant convolution, all heartbeat occurrences are identifiable as they are positively correlated with the kernel, minimizing deformations induced by the trend.

learning pipelines often perform baseline removal as a pre-processing step (Baek et al., 2015; Yan et al., 2019; Zhang et al., 2020b). However, these methods usually require extensive hyper-parameter tuning (Zhang et al., 2020a; Singhal et al., 2020). In deep learning, this challenge has been tackled through trend modeling components (Oreshkin et al., 2019), which has also been combined with contrastive learning (Liu et al., 2024; Woo et al., 2022a). Interestingly, convolution-based networks can leverage hard-coded invariant filters to accurately approximate trend invariance. Indeed, convolutions focus on local information, and similarly to spline functions, the trend can locally be approximated by low-degree polynomial functions, which, in its simplest form, can be seen as linear (1-degree). Figure 1 illustrates the cross-correlation between a single heartbeat and an ECG affected by a trend assuming a standard convolution and a convolution invariant to a linear trend. Surprisingly, the standard convolution fails to detect the correlation between the query heartbeat and the individual ones. On the contrary, the linear trend invariant convolution successfully identifies correlations to all underlying heartbeats and potentially offers more robustness for any ECG diagnosis (Liu et al., 2021b).

Following the previous observations, our work tackles the challenge of hard-coded spatiotemporal invariance within convolutional layers for time series accounting for deformations like offset shift or linear trend and extending beyond mere time invariance. In the literature, very few works deal with hard-coded time invariance, and most are concerned with local time-warping invariance (Shulman, 2019) or time rescaling invariance (Jacques et al., 2022). In addition, our mathematical framework provides an exact formulation for invariant convolutions, as opposed to approximating invariance in previous works. More specifically:

Section 3. We formulate time series deformations via group actions and introduce invariance under these actions. We then design generic and exact hard-coded invariant convolutions capable of handling deformations such as trends.

Section 4.1 We highlight the sensitivity of standard convolution to common deformations and demonstrate how deformation-invariant convolutions mitigate this issue, achieving better generalization than learned invariance.

Section 4.2, 4.3 & Appendix A.5.1 We showcase the effectiveness of the proposed convolutions against state-of-the-art baselines across multiple tasks. The results demonstrate that our hard-coded invariant convolutions offer a fast, generalizable, and robust approach to time series representation learning.

2. Related Work

Deep Learning for Time Series. Dominant deep learning frameworks for time series leverage multi-layer perceptrons (MLPs) (Oreshkin et al., 2019), convolutional networks (CNNs) (Bai et al., 2018) and recurrent ones (RNNs) (Salinas et al., 2020), as well as transformer-based networks (i.e., built upon the attention mechanism) (Wen et al., 2022). Convolutional kernels have traditionally dominated feature extraction in time series, from shapelets (Ye & Keogh, 2011) to the recently successful ROCKET (Dempster et al., 2020), that exploits several random kernels. Additionally, convolutional layers of different kernel sizes, often stacked in deep architectures (Ismail Fawaz et al., 2019), with increased receptive fields, such as INCEPTIONTIME (Ismail Fawaz et al., 2020) and RESNET (Wang et al., 2017), are prominent for time series classification. Similarly, TIMESNET (Wu et al., 2022) model capitalizes on convolutional layers to capture variations of multiple periodicities of 2D transformed multivariate time series, to solve multiple time series tasks. Beyond standard CNNs for time series, T-WaveNet (Minhao et al., 2021) is a tree-structured wavelet neural network that decomposes the input signal into various frequency subbands with similar energies based on the dominant frequency range. Another recent hierarchical CNN-based model for time series forecasting, SCINET (Liu et al., 2022), repeatedly downsamples and convolves

the input to enable information sharing at several resolutions. Furthermore, leveraging the success of the attention mechanism in text, transformer-based architectures have lately proven successful in capturing temporal interactions between multivariate time series inputs, mainly in time series forecasting (Zhou et al., 2021; 2022; Woo et al., 2022b; Liu et al., 2021a). For instance, AUTOFORMER (Wu et al., 2021) combines decomposition modules with an auto-correlation in place of self-attention, while CROSSFORMER (Zhang & Yan, 2022) capitalizes on 2D vector array embeddings that preserve temporal and channel information, followed by temporal and channel-wise cross-attention modules. Finally, several recent works, evaluate forecasting architectures also on the anomaly detection (Xu, 2021), by reformulating the task to point-wise reconstruction, with reconstruction error being the anomaly criterion. To overcome the sensitivity of transformers in overfitting, recent MLP-based architectures have showcased superior performance in forecasting, e.g., TSMIXER (Chen et al., 2023), FRETTS (Yi et al., 2024), while several studies doubt the robustness of the former in time series modeling (Zhang et al., 2022; Zeng et al., 2023).

Invariances for Time Series Modeling. Besides supervised methods, multiple approaches emphasize unsupervised learning for extracting representations from time series data before the downstream task (Nie et al., 2022; Dong et al., 2024). Among prominent approaches, self-supervised contrastive methods enforce invariances between representations by applying transformations (e.g., scaling, shifting, or noise injection) to the input and training the model to map them to the same underlying representations (Chen et al., 2020). Typically, CNNs constitute basic blocks for various time series augmentation-based contrastive frameworks, such as TS-TCC (Eldele et al., 2021) and TIMECLR (Yang et al., 2022). Except for transformation-based augmentations, samples can also be contrasted with sampled subsequences (Franceschi et al., 2019), adjacent segments (Tonekaboni et al., 2021), or a combination of all (Yue et al., 2022). Unlike deep learning, invariances in time series data have long been a central focus in classical time series data mining approaches (Esling & Agon, 2012). For instance, local time warping invariance (Ding et al., 2008) can be tackled by dynamic time warping (DTW). Traditionally, amplitude and offset invariances are accomplished by Z-normalizing the data (Paparrizos et al., 2020). Finally, the LT-normalized distance (Germain et al., 2024) extends Z-normalized distance by adding invariance to linear trends, improving similarity search and motif discovery.

3. Method

3.1. Invariant embedding for time series

We, next, develop a mathematical framework to create an embedding invariant to a predefined set of deformations.

Essentially, the embedding is expected to map a geometrical object or any of its deformed versions to the same representative. In addition, we present some important properties that should verify the embedding and tailor the proposed framework to the case of time series.

Deformations and group action. From a geometrical viewpoint, the notion of invariance depends on the representation of deformations and the definition of the action of a deformation on a geometrical object. A classical approach consists of representing a deformation as an element of a group and its action by a group action:

Definition 3.1 (Group action). A group G with neutral e acts on the left on a set M , if there exists a map $a : G \times M \mapsto M$ that verifies:

- 1) $a(e, m) = m, \quad \forall m \in M$
- 2) $a(g, a(h, m)) = a(gh, m), \quad \forall (g, h) \in G^2, \forall m \in M$.

To simplify notation, the left action of $g \in G$ on $m \in M$ is denoted $g \cdot m$. For a group G that acts on the left on a set M , the orbit of $m \in M$ is the set of all its deformed versions $[m] = \{g \cdot m \mid g \in G\}$. The set of independent orbits, denoted M/G , is called quotient space, and if this set is reduced to a singleton, the action of G on M is said transitive, and it verifies that for any $m \in M$ its orbit is the whole set: $[m] = M$.

A group action for time series. Leveraging measure theory, we model the set of time series by the Hilbert space $L^2(I, \mathbb{R}^D, \mu)$ of functions defined on the closed interval $I \subset \mathbb{R}$ taking value in \mathbb{R}^D and square-integrable for the Borel measure μ . The inner product on $L^2(I, \mathbb{R}^D, \mu)$ is defined as:

$$\langle f, g \rangle_L = \int_I \langle f(t), g(t) \rangle d\mu(t) \quad (1)$$

where $\langle \cdot, \cdot \rangle$ is the dot product on \mathbb{R}^D . Let H be a finite dimensional vector subspace of $L^2(I, \mathbb{R}^D, \mu)$, we model the group of deformations as the set $\mathbb{R}_+^* \ltimes H$ with the composition rule $(\lambda_2, h_2) \times (\lambda_1, h_1) = (\lambda_2 \lambda_1, h_2 + \lambda_2 h_1)$. Finally, we model the group action by the application:

$$\begin{aligned} (\mathbb{R}_+^* \ltimes H) \times L^2(I, \mathbb{R}^D, \mu) &\rightarrow L^2(I, \mathbb{R}^D, \mu) \\ ((\lambda, h), f) &\mapsto \lambda f + h \end{aligned} \quad (2)$$

This is a general group action that is not transitive as H is a finite-dimensional vector subspace of $L^2(I, \mathbb{R}^D, \mu)$. By convention, we refer to $\mathbb{R}_+^* \ltimes H$ as the set of rigid deformations. The customization of the group action depends on the choice of basis for the subspace H . For instance, the Z-normalization (Paparrizos et al., 2020) is an invariant offset shift which corresponds to the subspace of deformations $\{h : I \mapsto c \mid c \in \mathbb{R}^D\}$ with the basis $\{h_i : I \mapsto e_i / \sqrt{\text{length}(I)} \mid i \in [1, \dots, D]\}$ where $(e_i)_{i \in [1, \dots, D]}$ is the orthonormal basis of \mathbb{R}^D .

165 **Invariant embedding.** An embedding invariant to a group
 166 action is expected to map any element of an orbit to the same
 167 representative, and it is defined as follows:

168 **Definition 3.2** (Invariant & orbit-injective embedding). An
 169 embedding map $L : M \mapsto N$ is said to be G -invariant, if
 170 for any $(g, m) \in G \times M$, $L(g \cdot m) = L(m)$. Additionally,
 171 L is said to be orbit-injective if the application $\tilde{L} : [m] \in M/G \mapsto L(m) \in N$ is injective.
 172

173 Note that an invariant embedding is meaningful in the case
 174 of a non-transitive group action. In addition, if the embed-
 175 ding is orbit-injective, each orbit has a distinct representa-
 176 tive.

177 For now, we focus on the action of the finite-dimensional
 178 subspace H of a Hilbert space M by the usual vector addition:
 179 $(h, m) \in H \times M \mapsto m + h \in M$. The following proposition
 180 exhibits a H -invariant embedding that is also orbit-injective.

181 *Proposition 1.* Let P_H be the orthogonal projector on H , and
 182 I_d be the identity map on M , the embedding, $L = I_d - P_H$
 183 (the projector on H^\perp) is H -invariant and orbit-injective.
 184

185 *Proof.* See Appendix A.1. \square

186 *Remark 3.3.* If $(h_i)_{i \in \llbracket 1, N \rrbracket}$ is an orthonormal basis of the
 187 finite dimensional vector subspace H , then the orthogonal
 188 projector on H as an explicit formulation $P_H : m \in M \mapsto$
 189 $\sum_{i=1}^N \langle m, h_i \rangle_L h_i \in H$.

190 Invariance to amplitude scaling can easily be incorporated
 191 in an embedding defined by the previous proposition:

192 *Proposition 2.* Let $L : M \mapsto M$ be the H -invariant and
 193 orbit-injective embedding map induced by the orthogonal
 194 projector on H as defined in proposition 1. The embedding
 195 map:

$$\hat{L} : m \in M \mapsto \begin{cases} L(m)/\|L(m)\|_M & \text{if } m \in M \setminus H \\ 0_M & \text{else} \end{cases} \quad (3)$$

196 is $(\mathbb{R}_+^* \ltimes H)$ -invariant and orbit-injective.

197 *Proof.* $(\mathbb{R}_+^* \ltimes H)$ -invariance is due to the linearity and H -
 198 invariance of L , and the orbit-injectivity is induced by the
 199 linearity and orbit-injectivity of L . \square

200 **An example: the univariate Z-normalization.** We are
 201 looking for an embedding invariant to amplitude scale and
 202 offset shift in the case of univariate discrete time series.
 203 The set of time series is modeled by $L^2([0, l], \mathbb{R}, \mu)$ where
 204 $l \in \mathbb{N}^*$, $\mu = \sum_{i=1}^l \delta_i$ and δ_i is the dirac measure at i . The
 205 set offset shifts is the subspace generated by the unit norm
 206 function $e : t \in [0, l] \mapsto 1/\sqrt{l} \in \mathbb{R}$. According to Propo-
 207 sition 2 the invariant embedding of a non-constant function
 208 f is the function: $(f - \langle f, e \rangle_{L^2})/\|f - \langle f, e \rangle_{L^2}\|_L$ which

209 leads to $(f(i) - \mu_f)/(\sqrt{l}\sigma_f)$ where $\mu_f = l^{-1} \sum_{i=1}^l f(i)$
 210 and $\sigma_f^2 = l^{-1} \sum_{i=1}^l (f(i) - \mu_f)^2$.

3.2. Invariant convolution

211 CNNs have been successful in many applications related
 212 to time series, essentially becoming a key building block
 213 of the latest deep neural networks. Their success comes
 214 from their ability to capture local information in long time
 215 series. However, CNNs remain sensitive to some deforma-
 216 tions like amplitude scaling or offset shifts (Mallat, 2016).
 217 In this section, we propose a novel convolution that is invari-
 218 ant to rigid deformations at a local scale while remaining
 219 computationally efficient.

220 **The formalism.** Let $L_{loc}^2(\mathbb{R}, \mathbb{R}^D, \mu)$ be the set of signals,
 221 we assume that the signal in square integrable on any com-
 222 pact of \mathbb{R} . Let $L^2(I, \mathbb{R}^D, \mu)$ be the set of kernels where
 223 $I \subset \mathbb{R}$ is a closed interval. The classical convolution layer,
 224 named 1D-CNN, between a signal f and a kernel w is the
 225 signal:

$$f * w : u \in \mathbb{R} \mapsto \int_I \langle f(u+t), w(t) \rangle d\mu(t) \in \mathbb{R} \quad (4)$$

226 Let assume a group of rigid deformations G acting on
 227 $L^2(I, \mathbb{R}^D, \mu)$, and \hat{L} the G -invariant embedding map defined
 228 by Proposition 2. For any $u \in \mathbb{R}$, we can define the operator
 229 K_u^G that maps the restriction of any signal f on the closed
 230 interval $u + I$ to its G -invariant representative:

$$K_u^G : \begin{array}{ccc} L_{loc}^2(\mathbb{R}, \mathbb{R}^D, \mu) & \rightarrow & L^2(I, \mathbb{R}^D, \mu) \\ f & \mapsto & \hat{L}(t \in I \mapsto f(t+u)) \end{array} \quad (5)$$

231 Leveraging these operators we define the G -invariant convo-
 232 lution between a signal f and a kernel w as the signal:

$$f *^G w : u \in \mathbb{R} \mapsto \int_I \langle (K_u^G f)(t), w(t) \rangle d\mu(t) \in \mathbb{R} \quad (6)$$

233 **Fast computation.** For a group of rigid deformations
 234 $(\mathbb{R}_+^* \ltimes H)$ with $(h_i)_{i \in \llbracket 1, N \rrbracket}$ a basis of H , thanks to
 235 Remark 3.3, the inner product between the invariant
 236 representation of $f \in L^2(I, \mathbb{R}^D, \mu)$ and w can be
 237 decomposed as follows: $\langle \hat{L}(f), w \rangle_L = (\langle f, w \rangle_L -$
 238 $\sum_{i=1}^N \langle f, h_i \rangle_L \langle w, h_i \rangle_L) / \|L(f)\|_L$. Assuming discrete sig-
 239 nals, the computation of $\langle \hat{L}(f), w \rangle_L$ requires the compu-
 240 tation of $2N + 2$ dot products. However, convolving a
 241 batch of B signals of length L with the kernel, the number
 242 of inner products to compute drops from $BL(2N + 2)$ to
 243 $BL(N + 2) + N$ as the inner products between the kernel
 244 and the basis are shared across signals and subsequences.
 245 It leads to the time complexity $\mathcal{O}(BLNCW)$ where C is
 246 the number of channels, W is the kernel size and assuming

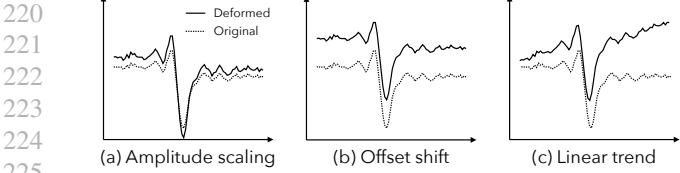


Figure 2: Deformations applied to an example series, including amplitude scaling, offset shift, and linear trend.

that $N \ll L$. Invariant convolutions do not consider small-size kernels (2 or 3 timestamps) but rather large kernels (30 or more). The traditional approach to convolution is not tractable in such a context. Instead, we leverage the Fast Fourier transform (FFT) (Mathieu et al., 2013), which changes the time complexity to $\mathcal{O}(BNCL \log(L))$. The computational time is identical for any window size, as the computation with the FFT does not depend on the kernel size. In the experimental results, we show that our proposed invariant convolutions benefit from fast computation.

4. Experimental Evaluation

We present an extended experimental evaluation illustrating the use and performance of hard-coded invariant convolutions. We specifically focus on convolutions invariant to constant functions (offset shift) $H_{\text{off}} = \{t \in I \mapsto b \mid b \in \mathbb{R}^D\}$ or affine functions (linear trend) $H_{\text{LT}} = \{t \in I \mapsto at + b \mid (a, b) \in \mathbb{R}^D \times \mathbb{R}^D\}$. As depicted in Introduction, these deformations are the simplest local approximation of trend deformation. In addition, we also include invariance to amplitude scaling, which is a common source of inter-individual variability. Figure 2 illustrates these three types of deformations using an ECG heartbeat example. The experimental evaluation is organized as follows:

I. Robustness to Deformations. We experimentally prove the robustness of the proposed invariant convolutions compared to vanilla ones and contrastive-based methods on a classification task when considering deformations.

II. Classification Performance. We show the competitive performance of an example architecture built upon a pool of normal and invariant convolutions on classification for several benchmark datasets. We also perform ablations and computational efficiency studies for the proposed method.

III. Generalization Performance. We assess the robustness of the proposed invariant convolutions on a transfer learning classification experiment against relevant contrastive learning methods.

IV. Anomaly Detection Performance. We finally capitalize on invariant filters to perform reconstruction-based anomaly detection by introducing an example decoder that effectively combines different types of features to recon-

struct the signal. The experimental protocol and the results are presented in Appendix A.5.1, revealing the benefit of invariant convolutions.

More information about the datasets and the implementation details can be found in Appendix A.3 and A.4, respectively. All implementations and results will be available on Github [SEE SUP. FILE].

4.1. Robustness to Deformations

Protocol. We aim to evaluate the robustness of trend deformations of hard-coded invariant convolutions (scaling/offset & scaling/linear trend) compared to vanilla convolutions and learned invariant representation with contrastive learning. Robustness is evaluated on a classification task. We consider 5 datasets from the UCR archive (Dau et al., 2019), which, by default, are all Z-normalized. Concerning baselines, we include three versions of inception-like (Ismail Fawaz et al., 2020) networks (a single convolution layer with kernels of four different sizes followed by a linear classifier) whose architecture is presented in Appendix A.2. The first one, CONV (normal), only includes standard convolutions, the second one, INV. CONV (offset), only includes amplitude/offset invariant convolutions, and the last one, INV. CONV (trend), only includes amplitude/linear trend invariant convolutions. These three models are only trained on the raw datasets and then tested on 5 different scenarios: (i) *no additional deformations*, (ii) *the addition of random offset*, sampled from uniform distribution between specific ranges, as well as (iii) *the addition of random trend* with slope and intercept values sampled again with uniform probability, (iv) *combination of added random offset and trend* and (v) *combination of random offset shift and smooth random walk*. For the last deformation, the added synthetic trend is a random walk generated from a Gaussian distribution and smoothed by a rolling mean. In order to compare with learned invariances, we also include the prominent supervised contrastive learning method TS-TCC (Eldele et al., 2021). On each dataset, this model is pre-trained with contrastive losses leveraging the different deformation scenarios to capture invariance; then, it is fine-tuned for classification on the raw data. Its performances are also evaluated under the 5 different scenarios.

Results. Table 1 displays accuracy scores, which are organized by scenarios of increasing magnitude of deformations, starting from (i) *no additional deformations* and ending with (v) *offset shift and smooth random walk*, the closer one to trend deformation.

In the first scenario (normalized data), there is no deformation-related distribution-shift shift between the train and the test sets. On average, CONV (normal) and TS-TCC perform slightly better than the hard-coded invariant net-

275 Table 1: Robustness study of distinct convolutional filter types with respect to invariance, i. e., standard ones (non-invariant),
 276 offset invariant, and trend invariant, on five *UCR* datasets. Models are trained on *normalized* data and tested on four
 277 additional synthetic deformation scenarios (random offset (off.), random linear trend (LT), smooth random walk (RW) and
 278 their combinations). Higher is better, best methods in **bold**, second best underlined. The self-supervised TS-TCC method is
 279 pre-trained with all deformation augmentations, and it is then fine-tuned (FT) on the raw normalized data. For the three
 280 convolutional variants, knowledge of the distribution of deformations is not incorporated during training, and thus, the
 281 model is evaluated in an out-of-distribution setting.

		CONV - normal -	INV. CONV - offset -	INV. CONV - trend -	TS-TCC (FT) (2021) - offset, trend -
HandOutlines	<i>Normalized</i>	77.84 ± 0.0	72.43 ± 0.54	70.54 ± 0.71	88.96 ± 1.17
	+ off.	41.08 ± 0.47	<u>72.43 ± 0.54</u>	<u>70.72 ± 1.02</u>	47.83 ± 0.54
	+ LT	37.39 ± 0.16	<u>71.08 ± 1.08</u>	<u>70.90 ± 0.82</u>	40.87 ± 1.41
	+ off., LT	35.68 ± 0.47	70.45 ± 1.22	<u>70.63 ± 0.41</u>	37.07 ± 0.87
	+ off., RW	35.95 ± 0.0	62.34 ± 1.02	<u>64.59 ± 1.24</u>	35.82 ± 0.15
MixedShapesRegularTrain	<i>Normalized</i>	93.91 ± 0.39	88.48 ± 1.79	92.26 ± 0.70	93.44 ± 1.51
	+ off.	38.96 ± 0.99	88.41 ± 1.83	<u>92.24 ± 0.73</u>	<u>92.56 ± 1.31</u>
	+ LT	29.33 ± 2.16	<u>84.63 ± 3.92</u>	<u>91.55 ± 0.69</u>	82.96 ± 4.72
	+ off., LT	28.22 ± 1.49	84.05 ± 4.08	<u>91.64 ± 0.80</u>	81.35 ± 4.08
	+ off., RW	22.94 ± 0.71	76.98 ± 4.62	<u>90.59 ± 0.64</u>	68.56 ± 4.86
NonInvasiveFetalECGThorax1	<i>Normalized</i>	91.35 ± 0.05	86.53 ± 0.06	85.00 ± 0.33	84.19 ± 1.18
	+ off.	15.32 ± 0.44	<u>85.11 ± 0.73</u>	<u>83.84 ± 0.12</u>	71.47 ± 6.51
	+ LT	8.09 ± 0.23	42.49 ± 0.33	<u>82.56 ± 0.45</u>	49.54 ± 5.97
	+ off., LT	6.09 ± 0.17	42.36 ± 0.65	<u>82.02 ± 0.52</u>	47.61 ± 6.70
	+ off., RW	3.48 ± 0.14	37.22 ± 0.94	<u>73.33 ± 0.05</u>	31.55 ± 6.16
FordB	<i>Normalized</i>	82.35 ± 0.57	<u>83.33 ± 0.5</u>	82.84 ± 0.65	78.37 ± 0.30
	+ off.	60.91 ± 3.17	<u>81.61 ± 1.38</u>	<u>83.37 ± 0.31</u>	78.29 ± 0.43
	+ LT	53.83 ± 1.10	70.50 ± 6.02	<u>82.68 ± 0.91</u>	76.85 ± 2.23
	+ off., LT	53.04 ± 0.75	69.47 ± 7.34	<u>82.35 ± 0.66</u>	76.49 ± 2.83
	+ off., RW	50.91 ± 0.26	68.80 ± 8.21	<u>80.37 ± 0.43</u>	75.77 ± 1.67
Yoga	<i>Normalized</i>	76.12 ± 0.24	76.80 ± 1.69	76.92 ± 0.14	<u>80.64 ± 0.63</u>
	+ off.	53.73 ± 0.26	72.07 ± 1.93	<u>75.71 ± 0.18</u>	74.97 ± 1.59
	+ LT	52.86 ± 0.14	67.71 ± 2.99	<u>74.17 ± 0.29</u>	62.21 ± 0.74
	+ off., LT	51.13 ± 0.07	64.85 ± 2.05	<u>72.88 ± 0.38</u>	61.85 ± 1.26
	+ off., RW	48.82 ± 0.12	65.67 ± 0.74	<u>72.58 ± 0.19</u>	62.69 ± 2.64
Percentage Drop (%) with respect to <u>Normalized</u> (lower the better in abs. value)	+ off.	-48% ± 20%	-1% ± 2%	0% ± 0%	-13% ± 17%
	+ LT	-55% ± 22%	-16% ± 17%	-1% ± 1%	-26% ± 19%
	+ off., LT	-57% ± 22%	-18% ± 17%	-1% ± 2%	-28% ± 20%
	+ off., RW	-59% ± 23%	-23% ± 16%	-6% ± 4%	-34% ± 22%

308 works INV. CONV (offset) and INV. CONV (trend). Outside
 309 of the *HandOutline* dataset, the performance difference be-
 310 tween the two groups does not exceed 5.0 points of accuracy.
 311 This slight performance difference could be attributed to
 312 small class amplitude, offset, or linear trend dependencies.
 313 However, as soon as deformations are added (even small),
 314 CONV (normal) and TS-TCC performances drop signif-
 315 icantly. In contrast, the performance of the deformation
 316 invariant network INV. CONV (trend) remains constant and
 317 becomes the best performer in most cases. On average, INV.
 318 CONV (trend) performs better than TS-TCC by 16.0 points
 319 of accuracy.

320 Regarding the specificities of each model, CONV (normal)
 321 is the most affected by deformations. Its performance
 322 drops by 48% when adding minor deformations (offset)
 323 and goes up to 59% with smooth random walk trends. It
 324 suggests that convolutions are highly sensitive to deforma-
 325 tions, even small ones, commonly observed in time series.
 326 Appendix A.5.2 clearly illustrates the sensitivity of CONV
 327 (normal) by comparing the feature maps of its convolutional
 328

329 filters with those of INV. CONV (offset) and INV. CONV
 330 (trend) on the same time series. When the time series un-
 331 dergoes deformation, the feature map landscape of CONV
 332 changes drastically. In contrast, its invariant counterparts
 333 (INV. CONV) maintain a consistent structure according
 334 to their invariance. TS-TCC is the second most affected
 335 model. On average, its performance drops by at most 34%,
 336 suggesting that the learned invariance is only partial and
 337 does not generalize well on new observations. In contrast,
 338 convolutions with hard-coded invariance behave accord-
 339 ing to their properties. INV. CONV (offset) performance
 340 remains constant when offset deformations are added but
 341 decreases when more complex deformations are added. In-
 342 terestingly, the performances are equivalent between the
 343 linear trend scenario and the offset + linear trend one, com-
 344 forting the invariance property to offset. INV. CONV (trend)
 345 performance remains almost constant on scenarios includ-
 346 ing offset and linear trend deformations ((ii),(iii),(iv)). Note
 347 that the slight performance variations in these three scenar-
 348 os are due to some padding effects. In the worst scenario,
 349 including a smooth random walk trend, the performance

330 Table 2: Classification Accuracy (%) for all considered datasets. Accuracy is averaged for datasets from *UEA* repository.
331 Higher is better, best methods in **bold**, second best underlined.

Datasets	INVCONVNET (ours)	TIMESNET (2022)	PATCHTST (2022)	CROSSFORMER (2022)	TSLANET (2024)	DLINEAR (2023)	INCEPTION (2020)	RESNET (2017)	CNN (2018)	ROCKET (2020)
UEA (26 datasets)	71.81 ± 0.80	<u>66.87 ± 1.72</u>	<u>66.18 ± 1.26</u>	<u>66.37 ± 1.35</u>	68.70 ± 1.19	61.51 ± 1.05	62.86 ± 1.96	67.37 ± 1.59	65.67 ± 1.64	71.29 ± 0.90
UCIHAR	96.63 ± 0.49	91.66 ± 0.62	85.74 ± 0.50	93.43 ± 0.56	94.71 ± 0.66	57.47 ± 0.73	95.26 ± 0.55	96.04 ± 0.48	95.78 ± 0.20	92.06 ± 0.15
Sleep-EDF	84.95 ± 0.39	74.64 ± 0.73	78.53 ± 0.28	79.82 ± 0.89	84.98 ± 0.43	36.15 ± 0.21	84.06 ± 0.39	85.62 ± 0.13	82.41 ± 0.56	83.88 ± 0.09
Epilepsy	98.43 ± 0.04	97.62 ± 0.20	98.01 ± 0.05	98.23 ± 0.12	98.23 ± 0.05	82.26 ± 0.06	97.65 ± 0.20	98.16 ± 0.04	97.61 ± 0.28	98.38 ± 0.02

338 Table 3: Classification ablation study of filter types
339 in INVCONVNET, including solely normal ones (-N),
340 offset-invariant (-O) ones, or trend-invariant ones (-T).
341

Datasets	INVCONVNET - mixed -	INVCONVNET-N - normal -	INVCONVNET-O - offset -	INVCONVNET-T - trend -
UEA (26 datasets)	71.81 ± 0.80	<u>68.04 ± 1.76</u>	67.70 ± 1.21	66.61 ± 2.58
UCIHAR	96.63 ± 0.49	96.13 ± 0.34	95.75 ± 0.51	95.43 ± 0.13
Sleep-EDF	84.95 ± 0.39	<u>84.70 ± 0.39</u>	83.73 ± 0.13	83.14 ± 0.13
Epilepsy	98.43 ± 0.04	98.09 ± 0.13	<u>98.26 ± 0.04</u>	<u>98.25 ± 0.09</u>

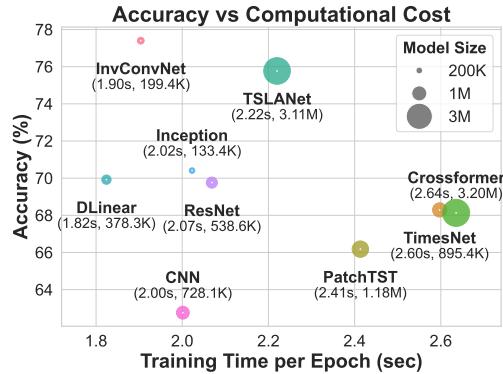
342 remains quite constant; it drops on average by 6% compared
343 to 59% for CONV (normal) and 34% for TS-TCC. This last
344 scenario is the closest to any trend deformation, indicating
345 that INV. CONV (trend) is well suited for classifying time
346 series affected by trends.

347 **Conclusion.** Standard convolutions are sensitive to spa-
348 tiotemporal deformations. This sensitivity issue can be
349 overcome by leveraging hard-coded invariant convolutions,
350 which also benefit from better generalization properties than
351 invariance learned by contrastive learning.

352 4.2. Classification Benchmark

353 **Pool of convolutions.** The choice of invariances is often
354 related to the application (Yue et al., 2022), and setting the
355 invariances by hand requires a good understanding of the
356 nature of the signals. In the absence of such knowledge,
357 as in classification benchmark, we decompose the space of
358 deformations $H = \bigoplus_{i=1}^K H_i$ in the direct sum of subspaces
359 such that the cumulative sums, $\emptyset \subset H_1 \subset H_1 + H_2 \subset \dots \subset H$,
360 represent sets of deformations of increasing order
361 of complexity. We consider a layer that concatenates of
362 n_j ($\bigoplus_{i=1}^j H_i$)-invariant convolutions for $j \in (1, \dots, K)$
363 and n_0 standard convolutions. In our specific experimental
364 framework, the linear trend deformations H_{LT} is the direct
365 sum of offset deformations $H_1 = \{t \in I \mapsto b \mid b \in \mathbb{R}^D\}$ and
366 purely linear deformations $H_2 = \{t \in I \mapsto at \mid a \in \mathbb{R}^D\}$. Il-
367 lustrations of the proposed pool of convolutions is provided
368 in the Figure 4 of Appendix A.2. Other pooling strategies,
369 like attention mechanism or reinforcement learning, are
370 possible and left for future work.

371 **Datasets.** We consider the 26 multivariate *UEA* datasets
372 (Bagnall et al., 2018), coming with a standard train/test
373 split. We also consider 3 additional datasets, the human



374 Figure 3: Models’ costs for *Heartbeat* dataset classification.

375 activity recognition *UCIHAR* (Anguita et al., 2013) dataset,
376 the *Sleep-EDF* dataset (Goldberger et al., 2000) for sleep
377 stage classification of EEG signals, and finally the epileptic
378 seizure recognition *Epilepsy* dataset (Andrzejak et al., 2001).
379 For these datasets, we follow the same preprocessing with
380 (Eldele et al., 2021), deriving train/validation/test sets of
381 60 : 20 : 20 ratio. Appendix A.3 provides additional details.

382 **Baselines.** We select nine state-of-the-art models for
383 time series classification. TIMESNET (Wu et al., 2022),
384 PATCHTST (Nie et al., 2022), CROSSFORMER (Zhang &
385 Yan, 2022) and DLINEAR (Zeng et al., 2023) are derived
386 from the Time-Series-Library (Wang et al., 2024). We also
387 compare to the CNN-based backbone of the self-supervised
388 method TSLANET (Eldele et al., 2024) and three powerful
389 CNN architectures, namely INCEPTION (Ismail Fawaz
390 et al., 2020), RESNET (Wang et al., 2017), and CNN (Ismail
391 Fawaz et al., 2018) build upon standard stacked 1D-CNNs.
392 We also include the state-of-the-art machine learning
393 method ROCKET (Dempster et al., 2020). Finally, our
394 model INVCONVNET combines a single inception-like pool
395 of convolutions layer of four different kernel sizes, combin-
396 ing balanced sets of standard, amplitude/offset and ampli-
397 tude/linear trend invariants convolutions. A linear classifier
398 follows the pool of convolutions. Implementation details
399 are presented in Appendix A.2.

400 **Results.** Table 2 presents the classification performance
401 of the proposed INVCONVNET against the nine considered

385 Table 4: Classification Accuracy (%) on a Transfer Learning experiment on *Fault-Diagnosis* (A, B, C, D sub-datasets) for
 386 the supervised INVCONVNET and INVCONVNET-N (normal) and two self-supervised methods.
 387

Methods	Train → Test												Avg. Acc. (%)
	A → B	A → C	A → D	B → A	B → C	B → D	C → A	C → B	C → D	D → A	D → B	D → C	
TS-TCC (FT)	55.33 ± 1.44	52.52 ± 4.55	62.13 ± 1.39	48.05 ± 3.32	71.50 ± 1.83	100.0 ± 0.0	40.76 ± 2.22	98.25 ± 1.22	99.34 ± 0.50	46.98 ± 0.65	100.0 ± 0.0	74.28 ± 2.77	70.76 ± 1.66
TS2Vec (FT)	54.11 ± 1.46	54.07 ± 1.91	52.54 ± 1.89	55.06 ± 0.17	88.72 ± 0.47	100.0 ± 0.0	57.81 ± 2.18	78.30 ± 3.80	78.41 ± 4.39	60.37 ± 1.95	99.97 ± 0.02	86.82 ± 0.54	72.18 ± 1.57
INVCONVNET (Sup.)	55.90 ± 0.42	55.93 ± 0.34	53.41 ± 0.14	85.10 ± 0.63	78.54 ± 0.17	99.05 ± 0.08	70.75 ± 1.32	85.04 ± 0.13	85.12 ± 0.15	70.91 ± 0.73	100.0 ± 0.0	78.49 ± 0.38	76.52 ± 0.37
INVCONVNET-N (Sup.)	60.55 ± 0.88	55.50 ± 1.82	53.50 ± 0.85	60.26 ± 2.01	77.30 ± 0.60	93.50 ± 0.90	64.93 ± 0.48	84.87 ± 0.23	84.46 ± 0.51	59.98 ± 0.98	99.96 ± 0.0	77.14 ± 0.14	72.66 ± 0.78

393 baselines, evaluated on *UEA* and the 3 additional datasets.
 394 All models are trained and tested for 3 runs with random
 395 seeds, and the average accuracy with its standard deviation
 396 is reported. Full classification results per dataset on
 397 *UEA* can be found in the Appendix A.5. We observe that
 398 on *UEA* repository, INVCONVNET has the best average
 399 test classification accuracy, followed closely by ROCKET,
 400 and both algorithms outperform the rest of deep learning
 401 approaches. ROCKET’s use of thousands of random convolu-
 402 tional kernels enables effective feature extraction, making
 403 it highly robust on smaller *UEA* classification datasets. On
 404 the additional 3 datasets, INVCONVNET shows superior
 405 performance in terms of accuracy, further validating the
 406 advantage of invariant CNN-based approaches in classifi-
 407 cation. For *Sleep-EDF*, RESNET is slightly better than the
 408 proposed INVCONVNET, which can be attributed to the
 409 depth of the method in capturing complex dependencies
 410 between the time series inputs. Finally, transformer-based
 411 methods are significantly outcompeted by CNN-based ones,
 412 and the MLP-based DLINERAR scores the worst, failing to
 413 capture the class-dependent temporal dynamics. Addition-
 414 ally, INVCONVNET built upon a single layer of convolutional
 415 kernels shows a significant advantage in terms of
 416 training time and memory cost, as presented in Figure 3 for
 417 the *UEA Heartbeat* dataset.

418
 419 **Ablation study.** In this experiment, we compare INVCONVNET classification performances to its standalone
 420 main components. More specifically, INVCONVNET com-
 421 bines balanced sets of standard, amplitude/offset, and
 422 amplitude/linear trend invariant convolutions. We com-
 423 pare this model to its standards INVCONVNET-N, ampli-
 424 tude/offset INVCONVNET-O, and amplitude/linear trend
 425 INVCONVNET-T counterparts. In all cases, the total num-
 426 ber of kernels remains the same. Table 3 classification
 427 performances of all four models. INVCONVNET shows
 428 the highest performance on all datasets. INVCONVNET-N
 429 is second in most cases, closely followed by the invariant
 430 models. It indicates that features invariants to amplitude,
 431 offset, or linear trend are valuable on many classification
 432 tasks while guaranteeing robustness to the networks as seen
 433 previously in Section 4.1, and by combining them, models
 434 achieve better performances.

435
 436 **Conclusion.** Convolutional features invariant to deforma-
 437 tions are meaningful on classification tasks, and their com-

392
 393 bination in single-layer offers simple and lightweight neural
 394 networks, offering comparable to better performances com-
 395 pared to the latest and prominent architectures.

4.3. Transfer Learning Experiment.

396 **Protocol.** We assess the generalization of invariant con-
 397 volutions on a transfer learning experiment using 4 dif-
 398 ferent source and target domains of the *Fault-Diagnosis*
 399 dataset (Lessmeier et al., 2016). Model wise, we include
 400 the self-supervised methods TS-TCC (Eldele et al., 2021)
 401 and TS2VEC (Yue et al., 2022) as well as INVCONVNET
 402 and its standard counterpart INVCONVNET-N configuration
 403 as presented in Section 4.2. Models are trained and tested
 404 on different source and target domains, namely the A, B, C,
 405 and D sub-datasets of *Fault-Diagnosis*, with direct transfer
 406 as in (Eldele et al., 2021). The self-supervised methods
 407 are pre-trained and fine-tuned (FT) on each source domain
 408 dataset leveraging contrastive learning.

409
 410 **Results.** Table 4 presents the transferred classification
 411 accuracy performances. Interestingly, the supervised INV-
 412 CONVNET, built on invariant convolutions, outperforms
 413 unsupervised methods and its standard counterpart (-N) by
 414 at least 4% with low variance.

415
 416 **Conclusion.** This experiment indicates that transfer learn-
 417 ing tasks on time series can benefit from convolutional in-
 418 variant features to remove the distribution shift caused by
 419 deformations and improve generalization.

5. Conclusion & Future Work

420 In this study, we establish a formal mathematical framework
 421 for time series invariances, and we design exact hard-coded
 422 invariant convolutions that seamlessly integrate within any
 423 CNN-based models. We experimentally show their en-
 424 hanced generalizability and computational efficiency in sev-
 425 eral setups. Exploring additional novel task-related invari-
 426 ances and associated merging strategies, along with unsu-
 427 pervised training, remains on our agenda for future work.

440 Impact Statement

441 The work presented in this paper aims to advance the field
442 of Machine Learning for time series analysis by introducing
443 a novel architectural design of convolutional layers that are
444 invariant to specific group actions for time series data, built
445 upon a thorough mathematical framework. Based on the pre-
446 sented experimental evaluation, the proposed invariant con-
447 volutional layers have the potential to significantly advance
448 diverse domains of time series modeling, such as healthcare,
449 environmental monitoring, and industrial machinery. By
450 enabling robust time series representations and insights into
451 their invariant properties, these layers can address critical
452 challenges in real-world applications and tasks, including
453 classification and anomaly detection, among others.

454 References

455 Abdulaal, A., Liu, Z., and Lancewicki, T. Practical ap-
456 proach to asynchronous multivariate time series anomaly
457 detection and localization. In *Proceedings of the 27th*
458 *ACM SIGKDD conference on knowledge discovery &*
459 *data mining*, pp. 2485–2494, 2021.

460 Andrzejak, R. G., Lehnertz, K., Mormann, F., Rieke, C.,
461 David, P., and Elger, C. E. Indications of nonlinear deter-
462 ministic and finite-dimensional structures in time series
463 of brain electrical activity: Dependence on recording re-
464 gion and brain state. *Physical Review E*, 64(6):061907,
465 2001.

466 Anguita, D., Ghio, A., Oneto, L., Parra, X., Reyes-Ortiz,
467 J. L., et al. A public domain dataset for human activity
468 recognition using smartphones. In *Esann*, volume 3, pp.
469 3, 2013.

470 Antoniou, A. Data augmentation generative adversarial
471 networks. *arXiv preprint arXiv:1711.04340*, 2017.

472 Baek, S.-J., Park, A., Ahn, Y.-J., and Choo, J. Baseline
473 correction using asymmetrically reweighted penalized
474 least squares smoothing. *Analyst*, 140(1):250–257, 2015.

475 Bagnall, A., Dau, H. A., Lines, J., Flynn, M., Large, J.,
476 Bostrom, A., Southam, P., and Keogh, E. The uea multi-
477 variate time series classification archive, 2018. *arXiv*
478 *preprint arXiv:1811.00075*, 2018.

479 Bai, S., Kolter, J. Z., and Koltun, V. An empirical evalua-
480 tion of generic convolutional and recurrent networks for
481 sequence modeling. *arXiv preprint arXiv:1803.01271*,
482 2018.

483 Biotti, A. and Mairal, J. Group invariance, stability to
484 deformations, and complexity of deep convolutional rep-
485 resentations. *Journal of Machine Learning Research*, 20
486 (25):1–49, 2019.

487 Box, G. E., Jenkins, G. M., Reinsel, G. C., and Ljung, G. M.
488 *Time series analysis: forecasting and control*. John Wiley
489 & Sons, 2015.

490 Bronstein, M. M., Bruna, J., Cohen, T., and Veličković,
491 P. Geometric deep learning: Grids, groups, graphs,
492 geodesics, and gauges. *arXiv preprint arXiv:2104.13478*,
493 2021.

494 Chen, S.-A., Li, C.-L., Yoder, N., Arik, S. O., and Pfis-
495 ter, T. Tsmixer: An all-mlp architecture for time series
496 forecasting. *arXiv preprint arXiv:2303.06053*, 2023.

497 Chen, T., Kornblith, S., Norouzi, M., and Hinton, G. A
498 simple framework for contrastive learning of visual rep-
499 resentations. In *International conference on machine
500 learning*, pp. 1597–1607. PMLR, 2020.

501 Dau, H. A., Bagnall, A., Kamgar, K., Yeh, C.-C. M., Zhu,
502 Y., Gharghabi, S., Ratanamahatana, C. A., and Keogh,
503 E. The ucr time series archive. *IEEE/CAA Journal of
504 Automatica Sinica*, 6(6):1293–1305, 2019.

505 Dempster, A., Petitjean, F., and Webb, G. I. Rocket: ex-
506 ceptionally fast and accurate time series classification
507 using random convolutional kernels. *Data Mining and
508 Knowledge Discovery*, 34(5):1454–1495, 2020.

509 Ding, H., Trajcevski, G., Scheuermann, P., Wang, X., and
510 Keogh, E. Querying and mining of time series data:
511 experimental comparison of representations and distance
512 measures. *Proceedings of the VLDB Endowment*, 1(2):
513 1542–1552, 2008.

514 Dong, J., Wu, H., Zhang, H., Zhang, L., Wang, J., and
515 Long, M. Simmtm: A simple pre-training framework
516 for masked time-series modeling. *Advances in Neural
517 Information Processing Systems*, 36, 2024.

518 Eldele, E., Ragab, M., Chen, Z., Wu, M., Kwoh, C. K., Li,
519 X., and Guan, C. Time-series representation learning
520 via temporal and contextual contrasting. *arXiv preprint
521 arXiv:2106.14112*, 2021.

522 Eldele, E., Ragab, M., Chen, Z., Wu, M., and Li, X. Tslanet:
523 Rethinking transformers for time series representation
524 learning. *arXiv preprint arXiv:2404.08472*, 2024.

525 Esling, P. and Agon, C. Time-series data mining. *ACM
526 Computing Surveys (CSUR)*, 45(1):1–34, 2012.

527 Franceschi, J.-Y., Dieuleveut, A., and Jaggi, M. Unsuper-
528 vised scalable representation learning for multivariate
529 time series. *Advances in neural information processing
530 systems*, 32, 2019.

531 Germain, T., Truong, C., and Oudre, L. Linear-trend nor-
532 malization for multivariate subsequence similarity search.

- 495 In 2024 IEEE 40th International Conference on Data
496 Engineering Workshops (ICDEW), pp. 167–175, 2024.
497 doi: 10.1109/ICDEW61823.2024.00028.
- 498 Goldberger, A. L., Amaral, L. A., Glass, L., Hausdorff, J. M.,
499 Ivanov, P. C., Mark, R. G., Mietus, J. E., Moody, G. B.,
500 Peng, C.-K., and Stanley, H. E. Physiobank, physiotoolkit,
501 and physionet: components of a new research resource
502 for complex physiologic signals. *Circulation*, 101(23):
503 e215–e220, 2000.
- 504 Hippke, M., David, T. J., Mulders, G. D., and Heller, R.
505 Wōtan: Comprehensive time-series detrending in python.
506 *The Astronomical Journal*, 158(4):143, 2019.
- 507 Horie, M., Morita, N., Hishinuma, T., Ihara, Y., and Mit-
508 sume, N. Isometric transformation invariant and equiv-
509 ariant graph convolutional networks. *arXiv preprint*
510 *arXiv:2005.06316*, 2020.
- 511 Hundman, K., Constantinou, V., Laporte, C., Colwell, I.,
512 and Soderstrom, T. Detecting spacecraft anomalies us-
513 ing lstms and nonparametric dynamic thresholding. In
514 *Proceedings of the 24th ACM SIGKDD international con-
515 ference on knowledge discovery & data mining*, pp. 387–
516 395, 2018.
- 517 Ismail Fawaz, H., Forestier, G., Weber, J., Idoumghar,
518 L., and Muller, P.-A. Evaluating surgical skills from
519 kinematic data using convolutional neural networks.
520 In *Medical Image Computing and Computer Assisted
521 Intervention—MICCAI 2018: 21st International Confer-
522 ence, Granada, Spain, September 16–20, 2018, Proceed-
523 ings, Part IV 11*, pp. 214–221. Springer, 2018.
- 524 Ismail Fawaz, H., Forestier, G., Weber, J., Idoumghar, L.,
525 and Muller, P.-A. Deep learning for time series classifi-
526 cation: a review. *Data mining and knowledge discovery*, 33
527 (4):917–963, 2019.
- 528 Ismail Fawaz, H., Lucas, B., Forestier, G., Pelletier, C.,
529 Schmidt, D. F., Weber, J., Webb, G. I., Idoumghar, L.,
530 Muller, P.-A., and Petitjean, F. Inceptiontime: Finding
531 alexnet for time series classification. *Data Mining and
532 Knowledge Discovery*, 34(6):1936–1962, 2020.
- 533 Jacques, B. G., Tiganj, Z., Sarkar, A., Howard, M., and
534 Sederberg, P. A deep convolutional neural network that
535 is invariant to time rescaling. In *International conference
536 on machine learning*, pp. 9729–9738. PMLR, 2022.
- 537 Keriven, N. and Peyré, G. Universal invariant and equivariant
538 graph neural networks. *Advances in Neural Information
539 Processing Systems*, 32, 2019.
- 540 Kitaev, N., Kaiser, Ł., and Levskaya, A. Reformer: The
541 efficient transformer. *arXiv preprint arXiv:2001.04451*,
542 2020.
- 543 Kvinge, H., Emerson, T., Jorgenson, G., Vasquez, S., Doster,
544 T., and Lew, J. In what ways are deep neural networks
545 invariant and how should we measure this? *Advances
546 in Neural Information Processing Systems*, 35:32816–
547 32829, 2022.
- 548 Lessmeier, C., Kimotho, J. K., Zimmer, D., and Sextro, W.
549 Condition monitoring of bearing damage in electromechanical
550 drive systems by using motor current signals of
551 electric motors: A benchmark data set for data-driven
552 classification. In *PHM Society European Conference*,
553 volume 3, 2016.
- 554 Liu, M., Zeng, A., Chen, M., Xu, Z., Lai, Q., Ma, L., and
555 Xu, Q. Scinet: Time series modeling and forecasting with
556 sample convolution and interaction. *Advances in Neural
557 Information Processing Systems*, 35:5816–5828, 2022.
- 558 Liu, S., Yu, H., Liao, C., Li, J., Lin, W., Liu, A. X., and
559 Dustdar, S. Pyraformer: Low-complexity pyramidal atten-
560 tion for long-range time series modeling and forecasting.
561 In *International conference on learning representations*,
562 2021a.
- 563 Liu, X., Wang, H., Li, Z., and Qin, L. Deep learning in
564 ecg diagnosis: A review. *Knowledge-Based Systems*, 227:
565 107187, 2021b.
- 566 Liu, Z., Alavi, A., Li, M., and Zhang, X. Guidelines for
567 augmentation selection in contrastive learning for time
568 series classification. *arXiv preprint arXiv:2407.09336*,
569 2024.
- 570 Löning, M., Bagnall, A., Ganesh, S., Kazakov, V., Lines,
571 J., and Király, F. J. sktime: A unified interface for
572 machine learning with time series. *arXiv preprint*
573 *arXiv:1909.07872*, 2019.
- 574 Mallat, S. Understanding deep convolutional networks.
575 *Philosophical Transactions of the Royal Society A: Mathe-
576 matical, Physical and Engineering Sciences*, 374(2065):
577 20150203, 2016.
- 578 Maron, H., Ben-Hamu, H., Shamir, N., and Lipman, Y.
579 Invariant and equivariant graph networks. *arXiv preprint*
580 *arXiv:1812.09902*, 2018.
- 581 Mathieu, M., Henaff, M., and LeCun, Y. Fast training
582 of convolutional networks through ffts. *arXiv preprint*
583 *arXiv:1312.5851*, 2013.
- 584 Mathur, A. P. and Tippenhauer, N. O. Swat: A water treat-
585 ment testbed for research and training on ics security. In
586 *2016 international workshop on cyber-physical systems
587 for smart water networks (CySWater)*, pp. 31–36. IEEE,
588 2016.

- 550 Middlehurst, M., Vickers, W., and Bagnall, A. Scalable
551 dictionary classifiers for time series classification. In
552 *Intelligent Data Engineering and Automated Learning–*
553 *IDEAL 2019: 20th International Conference, Manchester,*
554 *UK, November 14–16, 2019, Proceedings, Part I 20*, pp.
555 11–19. Springer, 2019.
- 556
- 557 Minhao, L., Zeng, A., Qiuxia, L., Gao, R., Li, M., Qin, J.,
558 and Xu, Q. T-wavenet: A tree-structured wavelet neural
559 network for time series signal analysis. In *International*
560 *Conference on Learning Representations*, 2021.
- 561
- 562 Moody, G. B. and Mark, R. G. The impact of the mit-bih
563 arrhythmia database. *IEEE engineering in medicine and*
564 *biology magazine*, 20(3):45–50, 2001.
- 565
- 566 Nie, Y., Nguyen, N. H., Sinthong, P., and Kalagnanam, J. A
567 time series is worth 64 words: Long-term forecasting with
568 transformers. *arXiv preprint arXiv:2211.14730*, 2022.
- 569
- 570 Oreshkin, B. N., Carpov, D., Chapados, N., and Bengio, Y.
571 N-beats: Neural basis expansion analysis for interpretable
572 time series forecasting. *arXiv preprint arXiv:1905.10437*,
573 2019.
- 574
- 575 Paparrizos, J., Liu, C., Elmore, A. J., and Franklin, M. J.
576 Debunking four long-standing misconceptions of time-
577 series distance measures. In *Proceedings of the 2020*
578 *ACM SIGMOD international conference on management*
579 *of data*, pp. 1887–1905, 2020.
- 580
- 581 Salinas, D., Flunkert, V., Gasthaus, J., and Januschowski,
582 T. Deepar: Probabilistic forecasting with autoregressive
583 recurrent networks. *International journal of forecasting*,
584 36(3):1181–1191, 2020.
- 585
- 586 Shulman, Y. Dynamic time warp convolutional networks.
587 *arXiv preprint arXiv:1911.01944*, 2019.
- 588
- 589 Singhal, A., Singh, P., Fatimah, B., and Pachori, R. B. An
590 efficient removal of power-line interference and baseline
591 wander from ecg signals by employing fourier decom-
592 position technique. *Biomedical Signal Processing and*
593 *Control*, 57:101741, 2020.
- 594
- 595 Su, Y., Zhao, Y., Niu, C., Liu, R., Sun, W., and Pei, D.
596 Robust anomaly detection for multivariate time series
597 through stochastic recurrent neural network. In *Proceed-
598 ings of the 25th ACM SIGKDD international conference*
599 *on knowledge discovery & data mining*, pp. 2828–2837,
2019.
- 600
- 601 Tonekaboni, S., Eytan, D., and Goldenberg, A. Unsuper-
602 vised representation learning for time series with temporal
603 neighborhood coding. *arXiv preprint arXiv:2106.00750*,
2021.
- 604
- Wang, Y., Wu, H., Dong, J., Liu, Y., Long, M., and Wang, J.
Deep time series models: A comprehensive survey and
benchmark. 2024.
- Wang, Z., Yan, W., and Oates, T. Time series classification
from scratch with deep neural networks: A strong base-
line. In *2017 International joint conference on neural*
networks (IJCNN), pp. 1578–1585. IEEE, 2017.
- Weiler, M. and Cesa, G. General e (2)-equivariant steerable
cnns. *Advances in neural information processing systems*,
32, 2019.
- Wen, Q., Zhou, T., Zhang, C., Chen, W., Ma, Z., Yan, J.,
and Sun, L. Transformers in time series: A survey. *arXiv*
preprint arXiv:2202.07125, 2022.
- Woo, G., Liu, C., Sahoo, D., Kumar, A., and Hoi, S. Cost:
Contrastive learning of disentangled seasonal-trend rep-
resentations for time series forecasting. *arXiv preprint*
arXiv:2202.01575, 2022a.
- Woo, G., Liu, C., Sahoo, D., Kumar, A., and Hoi, S. Ets-
former: Exponential smoothing transformers for time-
series forecasting. *arXiv preprint arXiv:2202.01381*,
2022b.
- Wu, H., Xu, J., Wang, J., and Long, M. Autoformer: Decom-
position transformers with auto-correlation for long-term
series forecasting. *Advances in neural information pro-
cessing systems*, 34:22419–22430, 2021.
- Wu, H., Hu, T., Liu, Y., Zhou, H., Wang, J., and Long,
M. Timesnet: Temporal 2d-variation modeling for gen-
eral time series analysis. In *The eleventh international*
conference on learning representations, 2022.
- Xu, J. Anomaly transformer: Time series anomaly de-
tection with association discrepancy. *arXiv preprint*
arXiv:2110.02642, 2021.
- Yan, H., Xu, T., Wang, P., Zhang, L., Hu, H., and Bai, Y.
Mems hydrophone signal denoising and baseline drift
removal algorithm based on parameter-optimized varia-
tional mode decomposition and correlation coefficient.
Sensors, 19(21):4622, 2019.
- Yang, X., Zhang, Z., and Cui, R. Timeclr: A self-supervised
contrastive learning framework for univariate time series
representation. *Knowledge-Based Systems*, 245:108606,
2022.
- Ye, L. and Keogh, E. Time series shapelets: a novel tech-
nique that allows accurate, interpretable and fast clas-
sification. *Data mining and knowledge discovery*, 22:
149–182, 2011.

- 605 Yi, K., Zhang, Q., Fan, W., Wang, S., Wang, P., He, H.,
606 An, N., Lian, D., Cao, L., and Niu, Z. Frequency-domain
607 mlps are more effective learners in time series forecasting.
608 *Advances in Neural Information Processing Systems*, 36,
609 2024.
- 610 Yue, Z., Wang, Y., Duan, J., Yang, T., Huang, C., Tong, Y.,
611 and Xu, B. Ts2vec: Towards universal representation of
612 time series. In *Proceedings of the AAAI Conference on*
613 *Artificial Intelligence*, volume 36, pp. 8980–8987, 2022.
- 614 Zaheer, M., Kottur, S., Ravanbakhsh, S., Poczos, B.,
615 Salakhutdinov, R. R., and Smola, A. J. Deep sets. *Advances in neural information processing systems*, 30,
616 2017.
- 617 Zeng, A., Chen, M., Zhang, L., and Xu, Q. Are transformers
618 effective for time series forecasting? In *Proceedings of*
619 *the AAAI conference on artificial intelligence*, volume 37,
620 pp. 11121–11128, 2023.
- 621 Zhang, F., Tang, X., Tong, A., Wang, B., and Wang, J.
622 An automatic baseline correction method based on the
623 penalized least squares method. *Sensors*, 20(7):2015,
624 2020a.
- 625 Zhang, F., Tang, X., Tong, A., Wang, B., Wang, J., Lv, Y.,
626 Tang, C., and Wang, J. Baseline correction for infrared
627 spectra using adaptive smoothness parameter penalized
628 least squares method. *Spectroscopy Letters*, 53(3):222–
629 233, 2020b.
- 630 Zhang, T., Zhang, Y., Cao, W., Bian, J., Yi, X., Zheng, S.,
631 and Li, J. Less is more: Fast multivariate time series
632 forecasting with light sampling-oriented mlp structures.
633 *arXiv preprint arXiv:2207.01186*, 2022.
- 634 Zhang, Y. and Yan, J. Crossformer: Transformer utilizing
635 cross-dimension dependency for multivariate time series
636 forecasting. In *The eleventh international conference on*
637 *learning representations*, 2022.
- 638 Zhang, Z.-M., Chen, S., and Liang, Y.-Z. Baseline correc-
639 tion using adaptive iteratively reweighted penalized least
640 squares. *Analyst*, 135(5):1138–1146, 2010.
- 641 Zhou, H., Zhang, S., Peng, J., Zhang, S., Li, J., Xiong, H.,
642 and Zhang, W. Informer: Beyond efficient transformer for
643 long sequence time-series forecasting. In *Proceedings of*
644 *the AAAI conference on artificial intelligence*, volume 35,
645 pp. 11106–11115, 2021.
- 646 Zhou, T., Ma, Z., Wen, Q., Wang, X., Sun, L., and Jin,
647 R. Fedformer: Frequency enhanced decomposed trans-
648 former for long-term series forecasting. In *Inter-
649 national conference on machine learning*, pp. 27268–27286.
650 PMLR, 2022.
- 651 Zhou, T., Niu, P., Sun, L., Jin, R., et al. One fits all: Power
652 general time series analysis by pretrained lm. *Advances in*
653 *neural information processing systems*, 36:43322–43355,
654 2023.

660 A. Appendix

661 A.1. Invariant embedding

663 Let M be a Hilbert space and H a finite dimensional vector subspace of M . We focus on the action H on M by the usual
664 vector addition: $(h, m) \in H \times M \mapsto m + h \in M$. The following proposition exhibits a H -invariant embedding that is also
665 orbit-injective.

666 *Proposition 3.* Let P_H be the orthogonal projector on H , and I_d be the identity map on M , the embedding, $L = I_d - P_H$ (the
667 projector on H^\perp) is H -invariant and orbit-injective.
668

669 *Proof. Existence of L :* As H is a finite dimension vector space, it is a closed and convex subset of the Hilbert space M ; the
670 orthogonal projector on H , denoted P_H , exists. Therefore, $L : m \in M \mapsto m - P_H(m) \in H$ is well defined.
671

672 **H-invariance of L :** Since H is closed, $M = H \oplus H^\perp$, and for any $x \in M$, we decompose $m = m_H + f_{H^\perp}$. Thus, for any
673 $m \in M$, and $h \in H$:

$$\begin{aligned} 674 L(m+h) &= m+h - P_H(m+h) \\ 675 &= m+h - P_H(m_{H^\perp} + m_H + h) \\ 676 &= m+h - (m_H + h) \quad (\textbf{projector on a closed vectorial subspace}) \\ 677 &= m - m_H \\ 678 &= L(m) \\ 679 \end{aligned}$$

680 which proves the H -invariance of L .

681 **Orbit-injectivity of L :** For any $m \in M$, its orbits corresponds to:

$$\begin{aligned} 683 [m] &= \{m+h \mid h \in H\} \\ 684 &= \{L(m) + h' \mid h \in H, h' = P_H(m) + h \in M\} \\ 685 &= L(m) + H \\ 686 \end{aligned}$$

687 Therefore, for any $([m], [m']) \in M/H \times M/H$, such that $[m] \cap [m'] = \emptyset$ implies that $L(m) \neq L(m')$ proving the
688 orbit-injectivity of L . \square
689

690 A.2. INVCONVNET: Architectural Details

691 The proposed pool of convolutions is presented in Figure 4 and is built upon the concatenation of normal filters, offset
692 shift-invariant filters, and linear trend invariant filters. The proposed convolutional layer, therefore, incorporates three
693 distinct kernel types with respect to invariance.
694

695 A.2.1. INVARIANT EMBEDDING MODULES

696 After mathematically formulating the characteristics of an invariant convolutional layer, which is built upon a standard (or
697 variant) kernel, a kernel invariant to offset shift and scaling, and a kernel invariant to linear trend and scaling, we provide
698 additional details for the design of the employed embedding modules. We present visualizations of the embedding modules
699 used for classification and anomaly detection (i. e., reconstruction) in Figure 5.
700

701 **Standard Module (Single-Layer):** The simplest embedding module is a single invariant convolutional layer for a specific
702 kernel size W and hidden dimensions d_{n_0} for the standard convolutional part (in yellow), d_{n_1} for the convolutional part
703 invariant to offset shift and scaling (in red), and d_{n_2} for the convolutional part invariant to linear trend and scaling (in
704 purple).
705

706 **Inception-like Module (Single-Layer):** We also study inception-like design by employing several kernel sizes, but without
707 stacking the layers in increasing depths. The depth of the employed module remains equal to one. As shown in Figure
708 5 (Left), in an inception-like embedding module, we consider several kernel sizes, e. g., W_1, W_2, W_3 , that are applied in
709 parallel to the input series, while leveraging the three parts of the proposed pool of convolutions (including standard and
710 invariant ones). The produced representation for the different kernel sizes, i. e., $\mathbf{z}_1, \mathbf{z}_2, \mathbf{z}_3$ are concatenated in the channel
711 dimension producing embeddings of size $(3 * \sum_{j=0}^2 d_{n_j}, L)$ for 3 selected kernel sizes. The distinct kernel sizes as well
712 as the hidden dimension for each part in the pool of convolutions, are hyperparameters that we need to tune, as in every
713 CNN-based architecture.
714

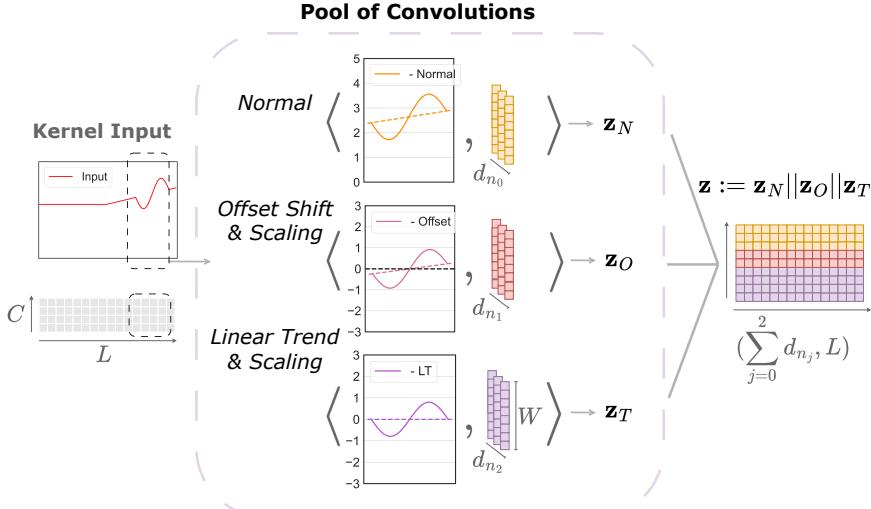


Figure 4: Visualization of the different kernel types employed on an input signal inside the proposed invariant convolutional layer, including normal filters (in yellow), filters invariant to offset shift (in red) and filters invariant to linear trend (in purple). The produced embedding z is the result of concatenations of the different representations.

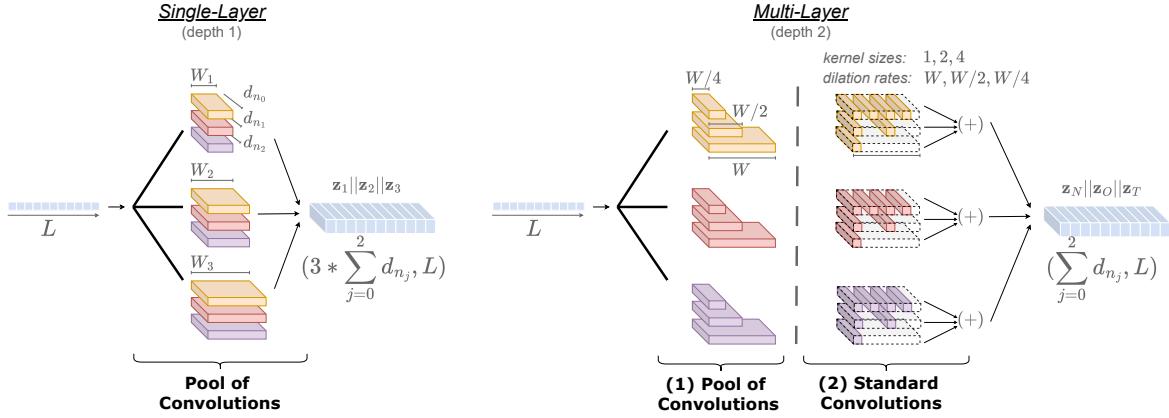
Multi-Scale Module (Multi-Layer): Additionally, we examine the capacity of a multi-scale embedding module built upon invariant convolutions as presented in Figure 5 (Right), particularly for the reconstruction task. Here the employed depth is equal to two. At the first level, an inception-like layer is employed, with kernel sizes selected to be powers of two (deriving the maximum exponent from the logarithm of half the series length and setting the minimum to four). The kernels, as described above, are applied in parallel, and the produced representation for the different kernel sizes is concatenated in the channel dimension. At the second level, a standard convolutional layer is applied. We similarly employ several kernel sizes of size $\max(W_i)/W_i$ matching the picked kernel sizes in the first layer W_i for $i \in \{1, \dots, K\}$, where K the number of kernels with distinct sizes. For each distinct kernel size in this second layer, a dilation factor is set as $r_i = W_i$, thus equal to the kernel size of the previous layer. This design enables capturing representations at different scales while employing a shallow and computationally light architecture, that still benefits from invariant convolutions. Experimentally, this module shows performance improvements for the anomaly detection task, where reconstruction can benefit from capturing dependencies at different granularities.

A.2.2. TASK-SPECIFIC MODULES

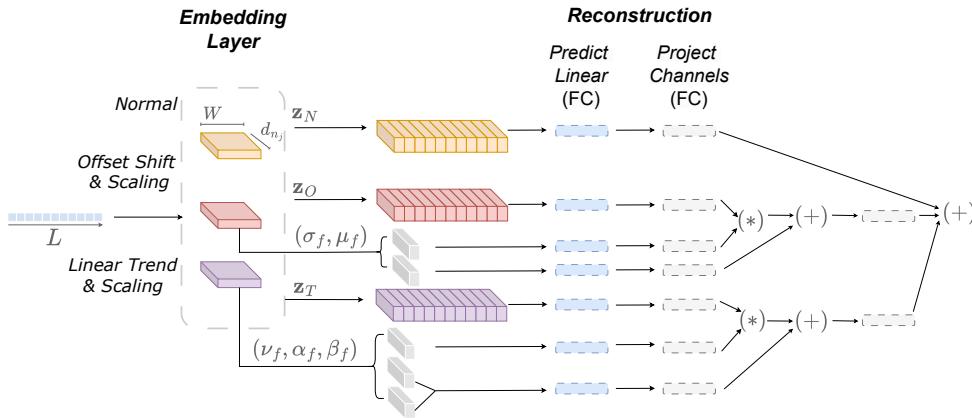
The embeddings derived from the modules described above are further processed by standard layers to produce the task-specific output.

I. Classification. The embedding is passed by a Global Average Pooling (GAP) layer, applied to the channel dimension, that averages over the time dimension to reduce the temporal features to a single value per channel. The result of the GAP layer is followed by a linear layer that produces the final class probabilities.

II. Anomaly Detection (Reconstruction). The multi-layer embedding module presented above is used to capture dependencies at various scales. For the reconstruction of the input, the embedding is followed by linear layers applied first on the temporal and then on the channel dimensions. The coefficients of the invariant kernels, i. e., (σ_f, μ_f) for the offset-invariant part and $(\nu_f, \alpha_f, \beta_f)$ for the trend-invariant part, are passed through linear layers to be mapped to the original time and channel dimensions and are then combined with each part (one variant and two invariant) of the embedding via addition and multiplication. More specifically, σ_f and ν_f refer to the norm of invariant embeddings of the signal, μ_f refers to the means, α_f and β_f refers to the coefficients of the linear trend. Combining the coefficients with the linearly projected embeddings adapts the level of the series to the original one, enabling enhanced temporal resolution. Figure 6 presents details about the reconstruction module and its relation with the embedding module.



788 Figure 5: **Left:** Single-layer embedding module, i.e., of depth 1, used in standard INVCONVNET (for 1 chosen kernel
789 width W) and in inception-like INVCONVNET (for several chosen kernel widths, e.g., 3 visualized in the figure). **Right:**
790 Multi-layer embedding module, i.e., of depth 2, used in multi-scale INVCONVNET (for several chosen kernel widths W).
791 In the second layer, for each kernel size, we utilize multiple kernels whose total number sums up to the larger kernel size,
792 achieving multi-scale views.



815 Figure 6: Visualization of the architecture used in terms of reconstruction that leverages the output of an embedding module
816 built upon a pool of convolutions (including variant and invariant ones). The representation of the embedding layer
817 z_N, z_O, z_T are passed from 2 fully connected linear layers (FC), from which the first operates on the temporal dimension
818 and the second on the channels. The coefficients of each invariant operation are similarly projected with linear layers and
819 combined with the representation (with addition or multiplication) to produce the output.

825 **A.3. Datasets Details**

826 We focus our experimental evaluation on several real-world time series datasets, including univariate and multivariate inputs,
 827 with significant applications in healthcare and medical diagnosis, wearable technology, audio processing, and transportation,
 828 among others.

829
 830 **Classification Datasets.** Details about the 26 multivariate derived from *UEA* data repository (Bagnall et al., 2018), that are
 831 employed in terms of this study in the classification experiment are provided in Table 5. More specifically, for each dataset
 832 we mention the number of channels, the length of the multivariate series, as well as the number of classes and the number of
 833 instances in the predefined train and test sets.

834
 835 Table 5: Details of *UEA* datasets used for classification.

Dataset	#Train	#Test	#Channels	Length	#Classes
<i>ArticularyWordRecognition</i>	275	300	9	144	25
<i>AtrialFibrillation</i>	15	15	2	640	3
<i>BasicMotions</i>	40	40	6	100	4
<i>Cricket</i>	108	72	6	1197	12
<i>Epilepsy</i>	137	138	3	206	4
<i>EthanolConcentration</i>	261	263	3	1751	4
<i>FaceDetection</i>	5890	3524	144	62	2
<i>FingerMovements</i>	316	100	28	50	2
<i>HandMovementDirection</i>	320	147	10	400	4
<i>Handwriting</i>	150	850	3	152	26
<i>Heartbeat</i>	204	205	61	405	2
<i>InsectWingbeat</i>	30000	20000	200	78	10
<i>JapaneseVowels</i>	270	370	12	29	9
<i>Libras</i>	180	180	24	51	6
<i>LSST</i>	2459	2466	6	36	14
<i>MotorImagery</i>	278	100	64	3000	2
<i>NATOPS</i>	180	180	24	51	6
<i>PEMS-SF</i>	267	173	963	144	7
<i>PenDigits</i>	7494	3498	2	8	10
<i>PhonemeSpectra</i>	3315	3353	11	217	39
<i>RacketSports</i>	151	152	6	30	4
<i>SelfRegulationSCP1</i>	268	293	6	896	2
<i>SelfRegulationSCP2</i>	200	180	7	1152	2
<i>SpokenArabicDigits</i>	6599	2199	13	93	10
<i>StandWalkJump</i>	12	15	4	2500	3
<i>UWaveGestureLibrary</i>	120	320	3	315	8

863 Table 6: Details of additional datasets used for classification.

Dataset	#Train	#Test	#Channels	Length	#Classes
<i>UCIHAR</i>	7352	2947	9	128	6
<i>Sleep-EDF</i>	25612	8910	1	3000	5
<i>Epilepsy</i>	9200	2300	1	178	2
<i>Fault-Diagnosis</i>	8184	2728	1	5120	3

871 Table 6 contains the same details for the additional datasets used for classification, i. e., the *UCIHAR* (Anguita et al.,
 872 2013) dataset, the *Sleep-EDF* dataset (Goldberger et al., 2000), and the *Epilepsy* dataset (Andrzejak et al., 2001). More
 873 specifically, *UCIHAR* data were collected by 30 volunteers performing various activities, including laying, standing, sitting,
 874 walking, walking downstairs, and walking upstairs. Volunteers' records were captured by a waist smartphone, including
 875 distinct measurements connected to acceleration and velocity signals. The *Sleep-EDF* dataset from the PhysioBank database
 876 consists of PolySomnoGraphic sleep recordings containing EEG, among other measurements. We consider only the EEG
 877 signals following previous studies (Eldele et al., 2021) and performed sleep stage classification, including awake, rapid eye
 878 movement, and non-rapid eye movements. Finally, the *Epilepsy* dataset consists of EEG brain activity measurements for
 879

epileptic seizure classification. Following the preprocessing of (Eldele et al., 2021), we perform binary classification after merging classes referring to non-epileptic seizure. For the transfer learning experiment, we utilized the *Fault-Diagnosis* dataset (Lessmeier et al., 2016), as preprocessed in (Eldele et al., 2021), which comprises of measurements under 4 different working conditions, perceived as different domains, and are assigned to 3 classes, including a healthy and two fault classes.

We also conduct a synthetic experiment on 5 large datasets from *UCR* repository (Dau et al., 2019). The selected datasets were picked from the equal-length datasets of the relevant repository that combine a large number of samples in the training and test sets, along with a large series length. Datasets with train set size and series length less than 100 were not considered in our subset. Details about the selected *UCR* datasets are given in Table 7. Since many *UCR* datasets are already preprocessed using Z-normalization to achieve zero mean and unit variance, they are not ideal for demonstrating the impact of our invariant layers on classification performance. This is the primary reason for conducting a synthetic experiment on these datasets rather than a conventional one with the whole repository.

Table 7: Details of the subset of *UCR* datasets used for the synthetic classification experiment.

Dataset	#Train	#Test	#Channels	Length	#Classes
<i>HandOutlines</i>	1000	370	1	2709	2
<i>MixedShapesRegularTrain</i>	500	2425	1	1024	5
<i>NonInvasiveFetalECGThorax1</i>	1800	1965	1	750	42
<i>FordB</i>	3636	810	1	500	2
<i>Yoga</i>	300	3000	1	426	2

Anomaly Detection Datasets. Furthermore, we present in Table 8 the five employed anomaly detection datasets after preprocessing them on non-overlapping subsequences of length 100, also showing the number of channels and the size of the train, validation, and test splits. The *SMD* dataset (Su et al., 2019) consists of data related to server machines collected at an internet company, while the *MSL* and *SMAP* (Hundman et al., 2018) datasets comprise of telemetry data from spacecraft monitoring systems. The *SWaT* (Mathur & Tippenhauer, 2016) dataset is a collection of sensor data from the operations of a critical infrastructure system. Finally, the *PSM* (Abdulaal et al., 2021) dataset contains measurements from application server nodes on an internet website.

Table 8: Details of datasets used for anomaly detection.

Dataset	#Train	#Val	#Test	#Channels	Length
<i>SMD</i>	566724	141681	708420	38	100
<i>MSL</i>	44653	11664	73729	55	100
<i>SMAP</i>	108146	27037	427617	25	100
<i>SWaT</i>	396000	99000	449919	51	100
<i>PSM</i>	105984	26497	87841	25	100

Data Splits and Preprocessing. As mentioned already in the main paper, for the proposed method, we do not normalize the data using Z-normalization for *UEA* and the rest 4 datasets used in classification, while the datasets from *UCR* are used for the synthetic experiment are derived normalized by the data source. On the contrary, all data are normalized for classification and the baselines, as well as for anomaly detection and all considered models (including the proposed INVCONVNET). For the *UEA* datasets, we do validation on the whole training set since the test sets are, in several cases, quite large, and thus, a small subset of the train set picked for validation can be a misleading indicator of performance. For the rest classification datasets, we perform a split into train/validation/test sets with a 60 : 20 : 20 ratio, following (Eldele et al., 2021). Similarly, for the five anomaly detection datasets, we split into train/validation/test sets with a 70 : 10 : 20 ratio (Xu, 2021).

A.4. Implementation Details

All experiments presented in this study were conducted on an Nvidia Tesla V100 GPU, with 40 cores and 756 GB of memory. We utilized the Adam optimizer with a learning rate of lr = 0.001 for both classification and unsupervised anomaly detection tasks. We also adopted a linear cosine annealing learning rate scheduler for INVCONVNET in classification. More specifically, the scheduler started the warmup phase with a learning rate equal to 0.001, linearly increasing the learning rate over the first 10 epochs to 0.01. After the warmup, it gradually reduced the learning rate using a cosine annealing schedule, down to 0.0001 by the end of training. For anomaly detection and the rest methods, we utilized a learning rate scheduler of

935 0.5 decrease rate per epoch. To have better estimates for the generalization performance of all models and, most importantly,
936 our proposed shallow modules, we performed 3 runs with random seeds for all considered datasets and tasks. Additional
937 details for each task and the hyperparameters of the models are given below.

938 **- Classification Task:** We trained the models for 100 epochs for all *UEA* datasets, the 5 *UCR* datasets and the *Fault-Diagnosis* dataset.
939 We performed early stopping during training, after 20 epochs of no improvement in the validation
940 accuracy for all models and kept the configuration of weights that correspond to the best validation accuracy during training.
941 The standard cross entropy loss was optimized during training for classification. For the INVCONVNET model, we considered
942 the inception-like embedding module of Figure 5 (Left) for all datasets of *UEA* except for *Epilepsy*, *EthanolConcentration*,
943 *Heartbeat* that we selected the standard embedding module of single kernel size. Finally, for *FingerMovements*, *Handwriting*,
944 *Libras* and *SelfRegulationSCPI* datasets, we employed the multi-scale embedding layer of Figure 5 (Right). The type of the
945 embedding layer, as well as the hyperparameters for the convolutional layers, e. g., kernel size and hidden dimension, were
946 selected through random search and the best performance on the validation set. For the rest of the classification datasets,
947 i. e., the *UCIHAR*, the *Sleep-EDF* and the *Epilepsy* datasets, we trained all models for 300 epochs with 20 epochs patience
948 and considered the inception-like embedding layer, since it was performing better on the validation set.
949

950 **- Unsupervised Anomaly Detection Task:** We trained the models for 10 epochs and stopped training if no improvements
951 had been made in terms of validation loss for 3 epochs, saving the best model weights on the validation set. We optimized
952 the models using the mean squared error (MSE) between the real input sequences and the reconstructed ones. For all five
953 anomaly detection datasets, we used the multi-scale embedding layer of Figure 5 (Right), followed by the reconstruction
954 module built upon linear layers in Figure 6.

955 **- Hyperparameter Selection:** We next provide more information about the selection of the kernel sizes and hidden
956 dimensions for the different embedding modules tested in terms of INVCONVNET. For the standard pool of convolutions
957 with one specific kernel size W , we chose the kernel size as the minimum value between the value 50 and half of the length
958 of the time series. For the inception-like embedding module, we selected several kernel sizes, such as 51, 75, 101, and 125,
959 or factors of those values for which the length of the series is proportional. Finally, for the first layer (pool of convolutions)
960 of the multi-scale module, we computed the kernel sizes as powers of two, starting from 16 up to a maximum of 128,
961 based on the logarithmic scaling of half the series length. For all modules, we tested hidden dimensions sizes for the pool
962 of convolutions in {32, 64, 128, 256} doing a split that enabled almost equal contribution for the three parts, i. e., normal,
963 invariant to offset shift and scaling, and invariant to linear trend and scaling. For instance for total hidden size equal to 32
964 the different parts had (12, 10, 10) hidden dimensions respectively, for 64 the split became (24, 20, 20) and so on.
965

966 For the common CNN-based baselines INCEPTION, RESNET, CNN, we tuned the number of convolutional layers, the
967 kernel sizes, and the hidden size of each layer. We followed a random search for a value between 2 and 6 for the number
968 of blocks and {32, 64, 128, 256} for the hidden dimensions, whereas for the kernel sizes, we used those proposed in the
969 relevant papers (Ismail Fawaz et al., 2020; Wang et al., 2017; Ismail Fawaz et al., 2018). All baselines’ implementations are
970 derived from the Time-Series-Library (Wang et al., 2024), with the configurations mentioned in the respective papers, and
971 the main code resources for performing the different tasks, e. g., classification and anomaly detection were adopted. We also
972 used ROCKET (Dempster et al., 2020) from sktime Library (Löning et al., 2019), with 3000 random convolutional kernels.
973 Finally, for the transfer learning classification experiment, the self-supervised contrastive TS-TCC and TS2VEC methods
974 were trained with their default parameters for classification as proposed in the respective papers (Eldele et al., 2021; Yue
975 et al., 2022), for 50 epochs for each phase of pre-training and fine-tuning.
976

977 A.5. Additional Results

978 A.5.1. RECONSTRUCTION-BASED ANOMALY DETECTION BENCHMARK

980 Reconstruction-based anomaly detection involves training a model to learn a compact representation of normal data by
981 reconstructing the input. The reconstruction error acts as the anomaly criterion, indicating whether the time series does not
982 conform to the normal patterns based on a chosen threshold.
983

984 **Datasets.** For unsupervised anomaly detection, we deploy the following benchmark datasets; *SMD* (Su et al., 2019),
985 *MSL* and *SMAP* (Hundman et al., 2018), *SWaT* (Mathur & Tippenhauer, 2016) and *PSM* (Abdulaal et al., 2021). We
986 apply standard preprocessing to extract non-overlapping sub-sequences and split them into train/validation/test sets with a
987 70 : 10 : 20 ratio (Xu, 2021; Wu et al., 2022).
988

990 **Baselines.** Focusing on reconstruction, we include ten time series (regression) models, most from Time-Series-Library
991 (Wang et al., 2024), including TIMESNET (Wu et al., 2022), PATCHTST (Nie et al., 2022), ETSFORMER (Woo et al., 2022b),
992 FEDFORMER (Zhou et al., 2022), AUTOFORMER (Wu et al., 2021), PYRAFORMER (Liu et al., 2021a), INFORMER (Zhou
993 et al., 2021), REFORMER (Kitaev et al., 2020), LIGHTTS (Zhang et al., 2022), DLINEAR (Zeng et al., 2023), and the
994 TSLANET (Eldele et al., 2024) backbone. For the decoder, we capitalize on the learned slope and intercept values for
995 each of the two invariant embedding parts, i.e., the Offset and Linear Trend, to adjust the projected embedding back to the
996 original temporal dimensions along with a linear layer. The projection to the initial channel dimension is then obtained by
997 applying a second channel-wise linear layer. Details on the INVCONVNET embedding and reconstruction modules are also
998 given in Appendix A.2.

1000 **Results.** Table 9 shows F1-scores (%) for the proposed model and baselines across 5 anomaly detection datasets. INV
1001 CONVNET performs best on the SWaT dataset and ranks third in average performance across all datasets, slightly behind
1002 TIMESNET. The latter’s superior performance can be attributed to its refined CNN blocks, which capture multiple periodicities
1003 for finer granularity in reconstruction. TSLANET achieves the highest average F1-score, benefiting from Fourier blocks
1004 before CNN modules to capture both short- and long-term dependencies. Other models, like the MLP-based DLINEAR
1005 and transformer-based FEDFORMER, also show competitive results. Notably, INVCONVNET excels with a single layer of
1006 invariant convolutions, proving the effectiveness of leveraging multi-scale invariances in shallow architectures for anomaly
1007 detection.

1008 Table 9: Anomaly Detection results in terms of the F1-score (%) for all considered datasets. Higher is better, best methods
1009 in **bold**, second best underlined.

Datasets	INVCONVNET (ours)	TIMESNET (2022)	PATCHTST (2022)	TSLANET (2024)	ETSFORMER (2022b)	FEDFORMER (2022)	LIGHTTS (2022)	DLINEAR (2023)	AUTOFORMER (2021)	PYRAFORMER (2021a)	INFORMER (2021)	REFORMER (2020)
SMD	84.05 ± 0.16	84.61 ± 0.56	84.15 ± 0.48	84.33 ± 0.17	79.69 ± 0.69	71.11 ± 0.02	83.04 ± 0.49	83.56 ± 0.14	71.16 ± 0.02	71.36 ± 0.01	71.17 ± 0.03	71.22 ± 0.01
MSL	80.68 ± 0.01	80.33 ± 0.79	78.67 ± 0.04	74.65 ± 0.78	75.98 ± 0.54	82.06 ± 0.14	80.39 ± 0.06	81.92 ± 0.01	82.08 ± 0.04	81.00 ± 0.08	<u>82.02 ± 0.11</u>	81.52 ± 0.08
SMAP	68.29 ± 0.07	69.18 ± 0.21	68.84 ± 0.01	80.26 ± 0.05	67.45 ± 0.74	68.71 ± 0.01	67.47 ± 0.02	67.32 ± 0.01	75.28 ± 1.64	67.76 ± 0.13	68.74 ± 0.12	73.30 ± 0.15
SWaT	92.82 ± 0.19	92.71 ± 0.04	88.38 ± 1.11	91.65 ± 0.27	92.67 ± 0.06	79.18 ± 0.01	<u>92.75 ± 0.01</u>	92.66 ± 0.01	79.18 ± 0.01	80.91 ± 0.38	79.75 ± 0.74	79.17 ± 0.01
PSM	96.34 ± 0.01	96.85 ± 0.27	96.12 ± 0.01	96.20 ± 0.03	95.23 ± 0.03	89.44 ± 0.88	95.50 ± 0.02	<u>96.66 ± 0.01</u>	88.25 ± 0.01	93.66 ± 0.13	90.55 ± 0.05	90.74 ± 0.09
Avg. F1 (%)	84.44 ± 0.09	<u>84.74 ± 0.37</u>	83.23 ± 0.33	85.42 ± 0.26	82.20 ± 0.41	78.10 ± 0.21	83.83 ± 0.12	84.42 ± 0.04	79.19 ± 0.65	78.94 ± 0.15	78.45 ± 0.21	79.19 ± 0.07

1012 Finally, in Table 10, we perform additional comparisons, in terms of anomaly detection, including the INVCONVNET
1013 model and the standard CNN-based variants, namely INCEPTION, RESNET and CNN originally proposed for classification.
1014 All models have an identical reconstruction module, with the exception of INVCONVNET, which also includes the signal
1015 decomposition coefficients on the invariant basis.

1016 Table 10: Anomaly Detection results for INVCONVNET and vanilla CNN-based methods. Performance mentioned in terms
1017 of the F1-score (%). Higher is better, best methods in **bold**, second best underlined.

Datasets	INVCONVNET + predict linear, project channels	INCEPTION	RESNET	CNN
SMD	84.05 ± 0.16	71.48 ± 0.11	76.20 ± 0.52	<u>77.31 ± 0.91</u>
MSL	80.68 ± 0.01	81.68 ± 0.08	<u>81.25 ± 0.09</u>	79.96 ± 0.20
SMAP	68.29 ± 0.07	68.63 ± 0.16	67.24 ± 0.65	67.00 ± 0.07
SWaT	92.82 ± 0.19	<u>82.69 ± 0.70</u>	80.93 ± 0.04	80.24 ± 0.95
PSM	96.34 ± 0.01	92.02 ± 0.35	92.30 ± 0.82	<u>93.30 ± 0.58</u>
Avg. F1 (%)	84.44 ± 0.09	79.30 ± 0.28	<u>79.58 ± 0.42</u>	<u>79.56 ± 0.54</u>

1018 As observed, our proposed method consistently outperforms all evaluated CNN-based variants, suggesting the effectiveness
1019 of combining invariances with their related signal coefficients for reconstruction.

1020 **Conclusion.** Invariant convolutional features combined with their related coefficients appear to be a concise and suited
1021 representation of time series for unsupervised anomaly detection based on reconstruction.

1022 A.5.2. ROBUSTNESS STUDY - VISUALIZATION OF FEATURE MAPS

1023 In Figure 7, we provide visualizations of the feature maps produced from different filter types (among the ones introduced
1024 in the paper), which are later incorporated in INVCONVNET example architecture, including normal ones (CONVNET
1025 (normal)), filters invariant to offset shift (INVCONVNET (offset)) and filters invariant to linear trend (INVCONVNET (trend)).
1026 Specifically, we consider FordB dataset from the large UCR datasets, considered in the Robustness Study of Table 1.
1027

For the considered dataset, the convolutional filters presented below are trained on normalized raw data and tested on four additional synthetic deformation scenarios: (i) *random offset* (*off.*), (ii) *random linear trend* (*LT*), (iii) *combined offset and trend* (*off., LT*), and (iv) *combined offset and smooth random walk* (*off., RW*). For the last deformation, the added synthetic trend is a random walk generated from a Gaussian distribution and smoothed by a rolling mean.

Extraction of Feature Maps. After passing the input series through the convolutional layer of each model and the activation function (i.e., $\text{ReLU}(\cdot)$), we extract the feature map for each filter (or each hidden dimension) corresponding to the largest considered kernel size. We recall that for the synthetic experiment on the 4 larger UCR datasets, we leveraged the inception-like embedding module of Figure 5 (Left), built upon 4 different kernel sizes with each having an equal total number of filters (or hidden dimensions equal to 128). Please note that averaging over all outputs produced by the layer for the several distinct kernel sizes produces multi-scale representations that produce similar (in terms of activated regions) but smoother maps (in terms of intensity values). The resulting feature maps, which are derived by the activated outputs for the largest kernel size, are essentially 2D representations with dimensions equal to the series length L and the 128 hidden dimensions.

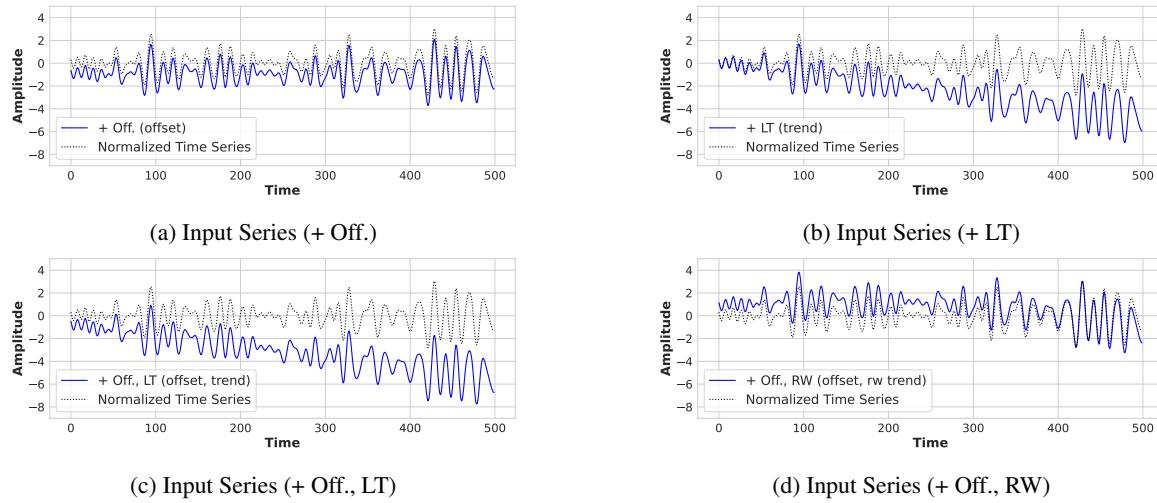


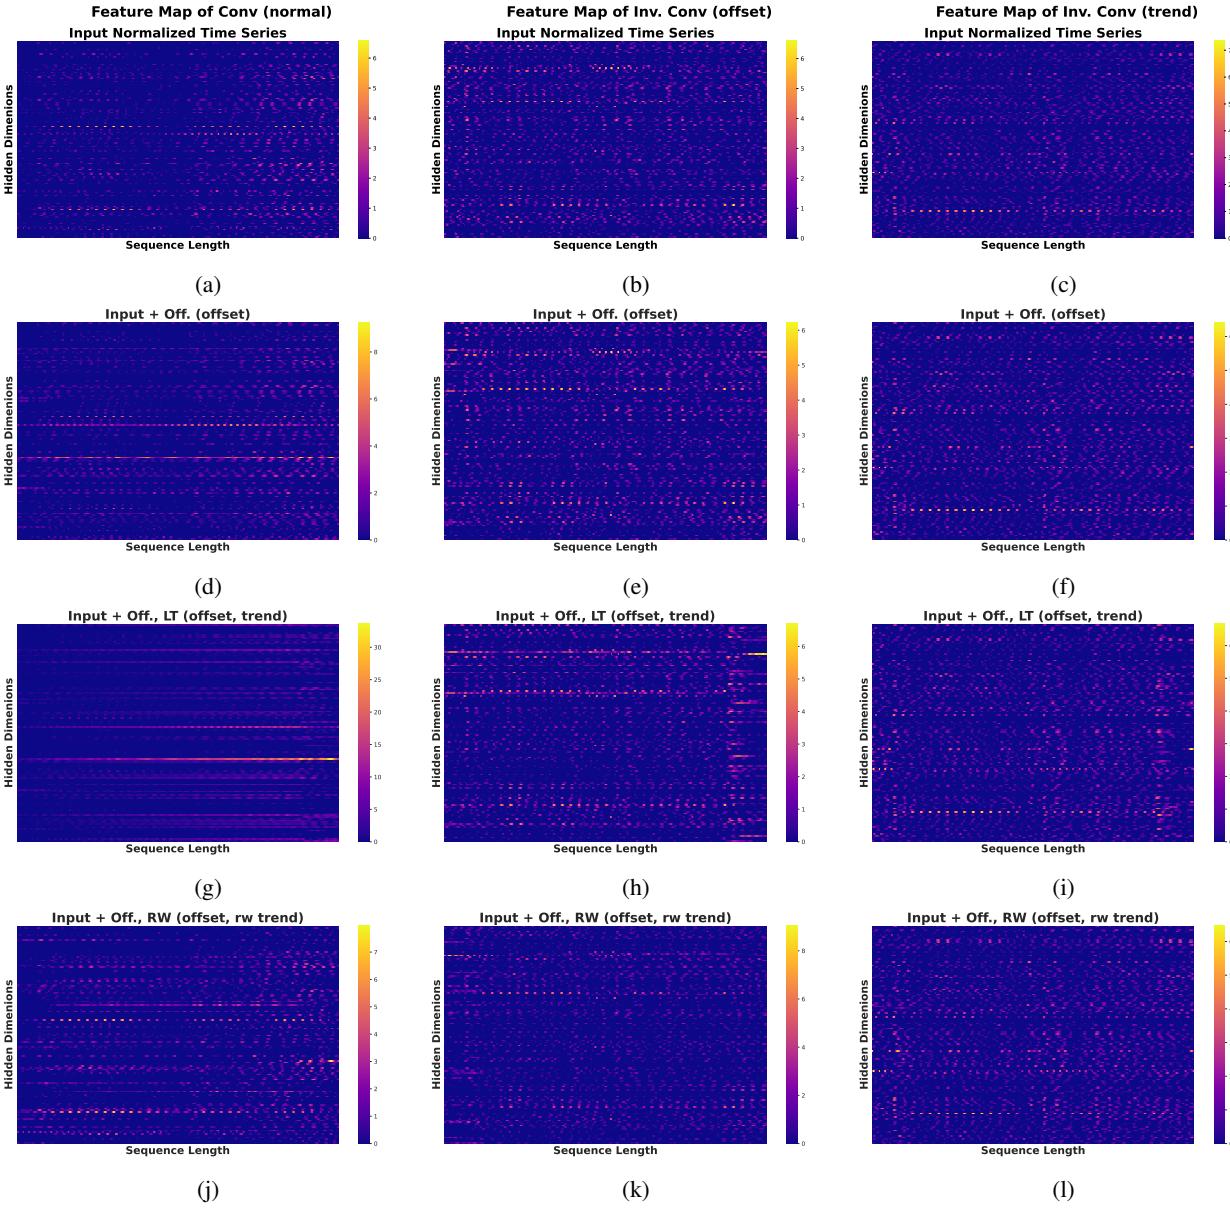
Figure 7: Plot of a single sample from the test set of *FordB* dataset for UCR used to produce the example feature maps of different types of convolutional filters. The normalized time series is captured in a black dotted line, while signals in blue represent the deformed versions of the series. Specifically, (a) represents the sample with the addition of random offset, (b) represents the sample with the addition of random linear trend, (c) represents the sample with the addition of random offset and linear trend, and (d) represents the sample with the addition of random offset and random walk trend.

We provide in Figure 8 the activated feature maps for a single sample for *FordB* dataset and its deformed versions (as presented in Figure 7), extracted for the different considered kernel types (CONV (normal), INV. CONV (offset) and INV. CONV (linear) as heatmaps. Color in the heatmap plots corresponds to the magnitude of the activation at a specific location of the series for each hidden dimension, with a lighter color (i.e., yellow) representing higher activations.

Deformation-Specific Activation Maps for Different Types of Filters. The feature maps of each considered convolutional filter type with respect to invariance, i.e., normal, offset shift-invariant, and trend-invariant, reveal distinct activation patterns under different deformation scenarios.

Normal filters are highly sensitive to both offset shifts and trends, often leading to widespread activation across the entire time series (See Figures 8d, 8g). This lack of selectivity causes the model to struggle to distinguish meaningful patterns from uninformative regions, which are better captured for the plain (normalized) data (Figure 8a). For instance, when a smooth random walk trend is introduced (Figure 8j), their activations become more uniform, indicating reduced sensitivity to meaningful variations in the input.

In contrast, offset shift-invariant filters remain robust to local shifts in the signal, exhibiting stable activation patterns even when an offset is introduced (See Figures 8b and 8e). However, their robustness is affected slightly when a linear trend is introduced (Figure 8h), as they begin to exhibit variability in their feature maps.



1143 Figure 8: Comparison of the feature maps produced for the robustness study (of Table 1) on a single sample of *FordB*
1144 dataset by the different types of convolutional filters before average-pooling for the normalized input and the 3 scenarios
1145 of synthetic deformations (addition of random offset, addition of random offset and linear trend and addition of random
1146 offset and random walk trend). Each column represents convolutional filters (128 in total) of the same type; (a),(d),(g),(j)
1147 correspond to normal convolutional filters (non-invariant), (b),(e),(h),(k) correspond to offset shift invariant convolutional
1148 filters and (b),(e),(h),(k) correspond to offset shift invariant convolutional filters and (c),(f),(i),(l).

Finally, trend-invariant filters consistently preserve their activation patterns across all tested deformations, demonstrating resilience to both offset shifts and trends (See all Figures 8f,8i,8l). This stability ensures that relevant features are captured effectively, regardless of input distortions. Similarly, trend-invariant filters are less affected by gradual changes, maintaining consistent responses across both plain and deformed data. While normal convolutions react strongly to local variations, offset and trend-invariant filters provide stability under deformations, offering a robust representation of the underlying time series patterns in the presence of deformations.

A.5.3. CLASSIFICATION

We also provide in Table 11 the full classification results for the 26 considered UEA datasets that correspond to the average of 3 runs for each combination of dataset and model. In the same table, we include again the already presented in the main paper, average accuracy for the whole collection of datasets as well as the number where each model scores first in the last row. The *JapaneseVowels* dataset is mentioned as out-of-time ('OOT') for not producing performance results since the experiment did not run within the time limits (12 hours maximum for each dataset). From the full classification results, we observe that the proposed INVCONVNET is, in several cases, slightly outperformed by the classical ROCKET method, but on average, is among the first best-competing models for most datasets, which explains its performance superiority in terms of average accuracy for the whole UEA.

Table 11: Full Classification results for UEA datasets. Accuracy (%) is mentioned for all combinations of models and datasets. Higher is better, best methods in **bold**, second best underlined.

Dataset	INVCONVNET	TIMESNET	PATCHTST	CROSSFORMER	TSLANET	DLINEAR	INCEPTION	RESNET	CNN	ROCKET
<i>ArticularWordRecognition</i>	99.00	97.78	97.67	98.22	98.22	96.67	84.56	98.44	97.89	99.44
<i>AtrialFibrillation</i>	<u>37.78</u>	28.89	42.22	28.89	24.44	35.56	28.89	24.44	33.33	6.67
<i>BasicMotions</i>	100.00	<u>95.00</u>	70.83	91.67	100.00	81.67	87.50	100.00	100.00	100.00
<i>Cricket</i>	98.61	93.06	94.44	92.59	97.69	91.20	87.96	98.15	<u>98.61</u>	100.00
<i>Epilepsy</i>	95.89	89.61	<u>97.34</u>	87.44	96.86	51.45	92.27	94.44	92.27	98.55
<i>EthanolConcentration</i>	25.98	26.24	23.32	39.67	22.18	24.97	23.57	21.93	22.81	<u>29.40</u>
<i>FaceDetection</i>	64.71	67.50	64.77	<u>65.26</u>	56.59	62.97	63.88	54.82	52.75	59.13
<i>FingerMovements</i>	56.33	<u>55.00</u>	53.33	52.33	<u>55.00</u>	48.67	56.33	53.00	53.00	54.00
<i>HandMovementDirection</i>	40.99	64.41	47.75	57.21	45.50	<u>59.01</u>	31.08	36.04	29.28	44.59
<i>Handwriting</i>	<u>53.14</u>	28.67	26.98	26.39	48.71	18.71	17.22	37.10	36.20	56.27
<i>Heartbeat</i>	77.40	68.29	66.18	68.13	<u>75.77</u>	69.92	70.41	69.76	62.76	73.17
<i>InsectWingbeat</i>	'OOT'	'OOT'	'OOT'	'OOT'	'OOT'	'OOT'	'OOT'	'OOT'	'OOT'	'OOT'
<i>JapaneseVowels</i>	97.66	91.71	94.68	96.76	96.85	93.33	91.80	98.83	98.38	97.39
<i>Libras</i>	88.70	79.07	76.11	86.30	84.81	50.19	57.04	94.07	88.70	<u>91.11</u>
<i>LSST</i>	55.04	12.77	48.35	11.21	10.41	31.85	35.71	8.99	9.37	60.76
<i>MotorImagery</i>	49.67	<u>52.00</u>	50.67	55.00	47.67	50.33	51.67	51.33	51.33	46.33
<i>NATOPS</i>	95.74	93.33	75.00	87.41	94.63	92.78	90.74	96.67	<u>95.74</u>	87.96
<i>PEMS-SF</i>	80.35	78.61	<u>81.89</u>	84.39	79.96	80.15	75.53	79.38	74.37	80.15
<i>PenDigits</i>	98.78	98.48	97.52	97.12	98.12	87.32	97.75	98.70	98.81	98.08
<i>PhonemeSpectra</i>	29.82	14.31	12.62	12.63	26.71	6.72	21.91	<u>28.66</u>	27.15	27.69
<i>RacketSports</i>	87.72	82.68	76.75	79.82	88.16	67.98	83.77	90.57	84.43	<u>90.35</u>
<i>SelfRegulationSCPI1</i>	<u>86.12</u>	87.60	78.84	85.32	79.18	83.39	81.57	80.32	84.64	84.53
<i>SelfRegulationSCPI2</i>	<u>54.44</u>	48.15	45.19	47.59	53.89	45.74	53.33	48.15	48.33	54.82
<i>SpokenArabicDigits</i>	99.47	98.83	97.92	98.67	99.58	95.85	98.50	99.23	98.98	<u>99.56</u>
<i>StandWalkJump</i>	28.89	33.33	51.11	24.44	<u>48.89</u>	33.33	26.67	40.00	31.11	<u>48.89</u>
<i>UWaveGestureLibrary</i>	<u>92.92</u>	86.35	83.02	84.69	87.71	77.92	61.77	81.35	71.56	93.33
Avg. Accuracy (%)	71.81	66.87	66.18	66.37	68.70	61.51	62.86	67.37	65.67	<u>71.29</u>
1st Count	4	3	2	3	2	0	1	5	2	8

In several studies (Wu et al., 2022; Zhou et al., 2023), only a subset of 10 UEA datasets is considered, and we also present once again the results for this subset along with total the average accuracy in 12. Similar observations can be made as those for Table 11, with the proposed INVCONVNET scoring the best average accuracy of 73.22%, followed by ROCKET.

1231 Table 12: Full Classification results for a subset of 10 UEA datasets. Accuracy (%) is mentioned for all combinations of
 1232 models and datasets. Higher is better, best methods in **bold**, second best underlined.

Dataset	INVCONVNET	TIMESNET	PATCHTST	CROSSFORMER	TSLANET	DLINEAR	INCEPTION	RESNET	CNN	ROCKET
<i>EthanolConcentration</i>	25.98	26.24	23.32	39.67	22.18	24.97	23.57	21.93	22.81	<u>29.40</u>
<i>FaceDetection</i>	64.71	67.50	64.77	<u>65.26</u>	56.59	62.97	63.88	54.82	52.75	59.13
<i>Handwriting</i>	<u>53.14</u>	28.67	26.98	26.39	<u>48.71</u>	18.71	17.22	37.10	36.20	56.27
<i>Heartbeat</i>	77.40	68.29	66.18	68.13	<u>75.77</u>	69.92	70.41	69.76	62.76	73.17
<i>JapaneseVowels</i>	97.66	91.71	94.68	96.76	96.85	93.33	91.80	98.83	<u>98.38</u>	97.39
<i>PEMS-SF</i>	80.35	78.61	<u>81.89</u>	84.39	79.96	80.15	75.53	79.38	74.37	80.15
<i>SelfRegulationSCP1</i>	<u>86.12</u>	87.60	78.84	85.32	79.18	83.39	81.57	80.32	84.64	84.53
<i>SelfRegulationSCP2</i>	<u>54.44</u>	48.15	45.19	47.59	53.89	45.74	53.33	48.15	48.33	54.82
<i>SpokenArabicDigits</i>	99.47	98.83	97.92	98.67	99.58	95.85	98.50	99.23	98.98	<u>99.56</u>
<i>UWaveGestureLibrary</i>	<u>92.92</u>	86.35	83.02	84.69	87.71	77.92	61.77	81.35	71.56	93.33
Avg. Accuracy (%)	73.22	68.20	66.28	69.69	70.04	65.30	63.76	67.09	65.08	<u>72.78</u>